

OSSANNA

UNIX PROGRAMMER'S MANUAL

K. Thompson

D. M. Ritchie

November 3, 1971

INTRODUCTION

This manual gives complete descriptions of all the publicly available features of UNIX. It provides neither a general overview (see "The UNIX Time-sharing System" for that) nor details of the implementation of the system (which remain to be disclosed).

Within the area it surveys, this manual attempts to be as complete and timely as possible. A conscious decision was made to describe each program in exactly the state it was in at the time its manual section was prepared. In particular, the desire to describe something as it should be, not as it is, was resisted. Inevitably, this means that many sections will soon be out of date. (The rate of change of the system is so great that a dismaying large number of early sections had to be modified while the rest were being written. The unbounded effort required to stay up-to-date is best indicated by the fact that several of the programs described were written specifically to aid in preparation of this manual!)

This manual is divided into seven sections:

- I. Commands
- II. System calls
- III. Subroutines
- IV. Special files
- V. File formats
- VI. User-maintained programs
- VII. Miscellaneous

Commands are programs intended to be invoked directly by the user, in contradistinction to subroutines, which are intended to be called by the user's programs. Commands generally reside in directory /bin (for binary programs). This directory is searched automatically by the command line interpreter. Some programs classified as commands are located elsewhere; this fact is indicated in the appropriate sections.

System calls are entries into the UNIX supervisor. In assembly language, they are coded with the use of the opcode "sys", a synonym for the trap instruction.

The special files section discusses the characteristics of each system "file" which actually refers to an I/O device.

The file formats section documents the structure of particular kinds of files; for example, the form of the output of the loader and assembler is given. Excluded are files used by only one command, for example the assembler's intermediate files.

User-maintained programs are not considered part of the UNIX system, and the principal reason for listing them is to indicate their existence without necessarily giving a complete

description. The author should be consulted for information.

The miscellaneous section gathers odds and ends.

Each section consists of a number of independent entries of a page or so each. The name of the entry is in the upper right corner of its pages, its preparation date in the upper left. Entries within each section are alphabetized. It was thought better to avoid page numbers, since it is hoped that the manual will be updated frequently.

All entries have a common format.

The name section repeats the entry name and gives a very short description of its purpose.

The synopsis summarizes the use of the program being described. A few conventions are used, particularly in the Commands section:

Underlined words are considered literals, and are typed just as they appear.

Square brackets ([]) around an argument indicate that the argument is optional. When an argument is given as "name", it always refers to a file name.

Ellipses "..." are used to show that the previous argument-prototype may be repeated.

A final convention is used by the commands themselves. An argument beginning with a minus sign "-" is often taken to mean some sort of flag argument even if it appears in a position where a file name could appear. Therefore, it is unwise to have files whose names begin with "-".

The description section discusses in detail the subject at hand.

The files section gives the names of files which are built into the program.

A see also section gives pointers to related information.

A diagnostics section discusses the diagnostics that may be produced. This section tends to be as terse as the diagnostics themselves.

The bugs section gives known bugs and sometimes deficiencies. Occasionally also the suggested fix is described.

The owner section gives the name of the person or persons to be consulted in case of difficulty. The rule has been that the last one to modify something owns it, so the owner is not necessarily the author. The owner's initials stand for:

ken	K. Thompson
dmr	D. M. Ritchie
jfo	J. F. Ossanna
rhm	R. Morris

These three-character names also happen to be UNIX user ID's, so messages may be transmitted by the mail command or, if the addressee is logged in, by write.

At the beginning of this document is a table of contents, organized by section and alphabetically within each section. There is also a permuted index derived from the table of contents. Within each index entry, the title of the writeup to which it refers is followed by the appropriate section number in parentheses. This fact is important because there is considerable name duplication among the sections, arising principally from commands which exist only to exercise a particular system call.

This manual was prepared using the UNIX text editor ed and the formatting program roff.

I. COMMANDS

ar	archive (combine) files
as	assembler
b	compile B program
bas	BASIC dialect
bcd	convert ASCII to BCD
boot	reboot system
cat	concatenate (or print) files
chdir	change working directory
check	check consistency of file system
chmod	change access mode of files
chown	change owner of files
cmp	compare file contents
cp	copy file
date	get date and time of day
db	symbolic debugger
dbppt	write binary paper tape
dc	desk calculator
df	find free disk space
dsw	delete files interactively
dtf	format DECtape
du	find disk usage
ed	text editor
find	find file with given name
for	compile Fortran program
form	generate form letter
hup	hang up typewriter
lbppt	read binary paper tape
ld	link editor (loader)
ln	link to file
ls	list contents of directory
mail	send mail to another user
mesg	permit or deny messages
mkdir	create directory
mkfs	initialize file system
mount	mount detachable file system
mv	move or rename file
nm	print namelist
od	octal dump of file
pr	print file with headings
rew	rewind DECtape
rkd	dump disk to tape
rkf	format RK disk
rkl	load disk from tape
rm	remove (delete) file
rmdir	remove (delete) directory
roff	run off (format) text
sdate	adjust date and time
sh	command interpreter
stat	get file status
strip	remove symbols, relocation bits
su	become super-user

sum	sum file
tap	manipulate DECtape
tm	get time information
tty	find name of terminal
type	print file on IBM 2741
umount	dismount removable file system
un	find undefined symbols
wc	get (English) word count
who	who is on the system
write	write to another user

II. SYSTEM CALLS

break	set program break
cemt	catch EMT traps
chdir	change working directory
chmod	change mode of file
chown	change owner of file
close	close open file
creat	create file
exec	execute program file
exit	terminate execution
fork	create new process
fstat	status of open file
getuid	get user ID
gtty	get typewriter mode
ilgins	catch illegal instruction trap
intr	catch or inhibit interrupts
link	link to file
mkdir	create directory
mount	mount file system
open	open file
quit	catch or inhibit quits
read	read file
rele	release processor
seek	move read or write pointer
setuid	set user ID
smdate	set date modified of file
stat	get file status
stime	set system time
stty	set mode of typewriter
tell	find read or write pointer
time	get time of year
umount	dismount file system
unlink	remove (delete) file
wait	wait for process
write	write file

III. SUBROUTINES

atof	convert ASCII to floating
atoi	convert ASCII to integer
ctime	convert time to ASCII
exp	exponential function

fptrap	floating-point simulator
ftoa	convert floating to ASCII
get	get character
itoa	convert integer to ASCII
log	logarithm base e
mesg	print string on typewriter
ptime	print time
putc	write character or word
sin	sine, cosine
switch	transfer depending on value

IV. SPECIAL FILES

mem	core memory as file
ppt	punched paper tape
rfo	RF disk file
rko	RK disk file
tap0,...,tap7	DECtape file
tty	console typewriter
tty0,...,tty5	remote typewriter

V. FILE FORMATS

a.out	assembler and loader output
archive	archive file
bppt	binary paper tape format
core	core image file
directory	directory format
file system	file system format
passwd	password file
uids	map names to user ID's
utmp	logged-in user information

VI. USER MAINTAINED PROGRAMS

basic	DEC supplied BASIC
bj	the game of black jack
cal	print calendar
chess	the game of chess
das	disassembler
dli	load DEC binary paper tapes
dpt	read DEC ASCII paper tapes
moo	the game of MOO
sort	sort a file
ttt	the game of tic-tac-toe

VII. MISCELLANEOUS

as2	assembler's pass 2
ascii	map of ASCII
ba	B assembler
bc	B compiler

bilib B interpreter library
bproc boot procedure
brt1,brt2 B start and finish
f1,f2,f3,f4 Fortran compiler passes
glob argument expander
init initializer process
kbd map of TTY 37 keyboard
liba standard assembly-language library
libb standard B library
libf standard Fortran library
login, logout logging on and logging off the system
msh mini Shell
suftab roff's suffix table
tabs set tab stops on typewriter

INDEX

chmod(I): change access mode of files
sdate(I): adjust date and time
mail(I): send mail to another user
write(I): write to another user

ar(I): a.out(V): assembler and loader output
archive(V): archive (combine) files
archive file
archive(V): archive file
argument expander
ar(I): archive (combine) files
ASCII paper tapes
bcd(I): convert ASCII to BCD
atof(III): convert ASCII to floating
atoi(III): convert ASCII to integer
ascii(VII): map of ASCII
ctime(III): convert time to ASCII
convert floating to ASCII...ftoa(III):
itoa(III): convert integer to ASCII

a.out(V): as(I): assembler and loader output
ba(VII): B assembler
as2(VII): assembler's pass 2 assembly-language library
liba(VII): standard as2(VII): assembler's pass 2
ba(VII): atof(III): convert ASCII to floating
bc(VII): atoi(III): convert ASCII to integer
B assembler
B compiler
B interpreter library
libb(VII): standard B library
b(I): compile B program
brt1,brt2(VII): B start and finish
log(III): logarithm base e
bas(I): bas(I): BASIC dialect
basic(VI): DEC supplied BASIC
BASIC
basic(VI): DEC supplied BASIC
ba(VII): B assembler
BCD
bcd(I): convert ASCII to BCD
bc(VII): B compiler
become super-user
b(I): compile B program
bilib(VII): B interpreter library
binary paper tape format
bppt(V): binary paper tape
dbppt(I): write binary paper tape
lbppt(I): read binary paper tape
dli(VI): load DEC binary paper tapes
remove symbols, relocation bits...strip(I):
bj(VI): the game of black jack

```

bj(VI): the game of black jack
bproc(VII): boot procedure
boot(I): reboot system
bppt(V): binary paper tape format
bproc(VII): boot procedure
break
break(II): set program break
brt1,brt2(VII): B start and finish
calculator
calendar
cal(VI): print calendar
catch EMT traps
catch illegal instruction trap
catch or inhibit interrupts
catch or inhibit quits
cat(I): concatenate (or print) files
cemt(II): catch EMT traps
change access mode of files
chmod(II): change mode of file
chown(II): change owner of file
chown(I): change owner of files
chdir(I): change working directory
chdir(II): change working directory
putc(III): write character or word
get(III): get character
check(I): check consistency of file system
chess(VI): the game of chess
chess(II): change working directory
check(I): check consistency of file system
chess
chess(VI): the game of chess
chmod(I): change access mode of files
chmod(II): change mode of file
chown(I): change owner of files
chown(II): change owner of file
close(II):
close open file
close(II): close open file
cmp(I): compare file contents
ar(I): archive
sh(I): command interpreter
cmp(I): compare file contents
b(I): compile B program
for(I): compile Fortran program
compiler passes
f1,f2,f3,f4(VII): Fortran
bc(VII): B
cat(I): concatenate (or print) files
check(I): check consistency of file system
tty(IV): console typewriter
ls(I): list contents of directory
contents
cmp(I): compare file
bcd(I): convert ASCII to BCD
atof(III): convert ASCII to floating
atoi(III): convert ASCII to integer
ftoa(III): convert floating to ASCII
itoa(III): convert integer to ASCII
ctime(III): convert time to ASCII

```

cp(I):	copy file
core(V):	core image file
mem(IV):	core memory as file
core(V):	core image file
cosine	
count	
cp(I):	copy file
mkdir(I):	create directory
mkdir(II):	create directory
creat(II):	create file
fork(II):	create new process
	creat(II): create file
	ctime(III): convert time to ASCII
	das(VI): disassembler
	date and time of day
	date and time
	date modified of file
	date(I): get date and time of day
	day
db(I):	symbolic debugger
dbpt(I):	write binary paper tape
dc(I):	desk calculator
debugger	
dpt(VI):	DEC ASCII paper tapes
dli(VI):	DEC binary paper tapes
basic(VI):	DEC supplied BASIC
tap0,...,tap7(IV):	DECtape file
dtf(I):	DECtape
rew(I):	DECtape
tap(I):	DECtape
rmdir(I):	(delete) directory
rm(I):	(delete) file
dsw(I):	delete files interactively
unlink(II):	(delete) file
mesg(I):	deny messages
switch(III):	depending on value
dc(I):	desk calculator
mount(I):	detachable file system
bas(I):	df(I): find free disk space
directory(V):	dialect
chdir(I):	directory format
chdir(II):	directory
ls(I):	directory
mkdir(I):	directory
mkdir(II):	directory
rmdir(I):	directory
	directory(V): directory format
	disassembler
	das(VI):
rf0(IV):	disk file
rk0(IV):	disk file
rkl(I):	disk from tape
df(I):	disk space
rkd(I):	disk to tape
du(I):	disk usage
rkf(I):	disk

umount(II):	dismount file system
umount(I):	dismount removable file system
dli(VI):	load DEC binary paper tapes
dpt(VI):	read DEC ASCII paper tapes
dsw(I):	delete files interactively
dtf(I):	format DECTape
du(I):	find disk usage
rkd(I):	dump disk to tape
od(I):	dump of file
ld(I):	link
ed(I):	text
log(III):	logarithm base
cemt(II):	catch
wc(I):	get
exec(II):	
exit(II):	terminate
glob(VII):	argument
exp(III):	
cmp(I):	compare
type(I):	print
stat(I):	get
stat(II):	get
file system(V):	
check consistency of	
mkfs(I):	initialize
mount(I):	mount detachable
mount(II):	mount
umount(I):	dismount removable
umount(II):	dismount
find(I):	find
pr(I):	print
archive(V):	archive
chmod(II):	change mode of
chown(II):	change owner of
close(II):	close open
core(V):	core image
cp(I):	copy
creat(II):	create
exec(II):	execute program
fstat(II):	status of open
link(II):	link to
ln(I):	link to
mem(IV):	core memory as
mv(I):	move or rename
od(I):	octal dump of
open(II):	open
passwd(V):	password
read(II):	read
rf0(IV):	RF disk
rk0(IV):	RK disk
rm(I):	remove (delete)

```

        dsw(I): delete files interactively
ar(I): archive (combine) files
    concatenate (or print) files...cat(I):
    change access mode of files...chmod(I):
chown(I): change owner of files
    set date modified of file...smdate(II):
        sort(VI): sort a file
            sum(I): sum file
tap0,...,tap7(IV): DECtape file
unlink(II): remove (delete) file
du(I): find disk usage
find(I): find file with given name
df(I): find free disk space
tty(I): find name of terminal
tell(II): find read or write pointer
un(I): find undefined symbols
find(I): find file with given name
finish
ftoa(III): convert floating to ASCII
atof(III): convert ASCII to floating
fptrap(III): floating-point simulator
form(I): generate for(I): compile Fortran program
rkf(I): format RK disk
roff(I): run off form letter
bppt(V): binary paper tape format
directory(V): directory format
file system(V): file system format
f1,f2,f3,f4(VII): form(I): generate form letter
libf(VII): standard Fortran compiler passes
for(I): compile Fortran library
df(I): find Fortran program
rkl(I): load disk fptrap(III): floating-point simulator
from tape fstat(II): status of open file
ftoa(III): convert floating to ASCII
exp(III): exponential function
bj(VI): the f1,f2,f3,f4(VII): Fortran compiler passes
chess(VI): the game of black jack
moo(VI): the game of chess
ttt(VI): the game of MOO
        form(I): game of tic-tac-toe
        get(III): generate form letter
        date(I): get character
        wc(I): get date and time of day
stat(I): get (English) word count
stat(II): get file status
        tm(I): get file status
time(II): get time information
gtty(II): get time of year
getuid(II): get typewriter mode
get(III): get user ID
get(III): get character

```

find(I): find file with
 hup(I):
 pr(I): print file with
 type(I): print file on
 getuid(II): get user
 setuid(II): set user
 uids(V): map names to user
 ilgins(II): catch illegal instruction trap
 core(V): core
 tm(I): get time
 utmp(V): logged-in user
 intr(II): catch or
 quit(II): catch or
 mkfs(I):
 init(VII):
 ilgins(II): catch illegal instruction trap
 itoa(III): convert integer to ASCII
 atoi(III): convert ASCII to integer
 dsw(I): delete files
 bilib(VII): B interpreter library
 sh(I): command interpreter
 intr(II): catch or inhibit
 bj(VI): the game of black
 kbd(VII): map of TTY 37
 form(I): generate form library...
 bilib(VII): B interpreter library
 standard assembly-language
 libb(VII): standard B library
 libf(VII): standard Fortran library
 ld(I): link editor (loader)
 link(II): link to file
 ln(I): link to file
 ls(I): list contents of directory
 dli(VI): load DEC binary paper tapes
 rkl(I): load disk from tape
 a.out(V): assembler and loader output
 ld(I): link editor (loader)
 log(III): logarithm base e
 utmp(V): logged-in user information
 logout(VII): logging on and off the system...login,
 getuid(II): get user ID given name
 glob(VII): argument expander
 gtty(II): get typewriter mode
 hang up typewriter
 headings
 hup(I): hang up typewriter
 IBM 2741
 ID
 ID
 ID's
 ilgins(II): catch illegal instruction trap
 image file
 information
 inhibit interrupts
 inhibit quits
 initialize file system
 initializer process
 init(VII): initializer process
 instruction trap
 integer
 interactively
 interpreter library
 interpreter
 interrupts
 intr(II): catch or inhibit interrupts
 itoa(III): convert integer to ASCII
 jack
 kbd(VII): map of TTY 37 keyboard
 keyboard
 lbppt(I): read binary paper tape
 ld(I): link editor (loader)
 letter
 liba(VII): standard assembly-language library
 libb(VII): standard B library
 libf(VII): standard Fortran library
 library
 library...liba(VII):
 library
 library
 link editor (loader)
 link to file
 link to file
 link(II): link to file
 list contents of directory
 ln(I): link to file
 load DEC binary paper tapes
 load disk from tape
 loader output
 (logarithm base e)
 logged-in user information
 logging off the system...login,

login, logout(VII):	logging on and logging off the system
logging off the system...	log(III): logarithm base e
the system...login,	login, logout(VII): logging on and
mail(I): send	logout(VII): logging on and logging off
tap(I):	ls(I): list contents of directory
uids(V):	mail to another user
ascii(VII):	mail(I): send mail to another user
kbd(VII):	manipulate DECtape
mem(IV): core	map names to user ID's
	map of ASCII
	map of TTY 37 keyboard
mesg(I): permit or deny	mem(IV): core memory as file
msh(VII):	memory as file
	mesg(I): permit or deny messages
	mesg(III): print string on typewriter
	messages
	mini Shell
chmod(II): change	mkdir(I): create directory
chmod(I): change access	mkdir(II): create directory
stty(II): set	mkfs(I): initialize file system
gtty(II): get typewriter	mode of file
smdate(II): set date	mode of files
moo(VI): the game of	mode of typewriter
	mode
mount(I):	modified of file
mount(II):	MOO
	moo(VI): the game of MOO
	mount detachable file system
	mount file system
	mount(I): mount detachable file system
	mount(II): mount file system
	move or rename file
mv(I):	move read or write pointer
seek(II):	msh(VII): mini Shell
	mv(I): move or rename file
	name of terminal
tty(I): find	name
find(I): find file with given	namelist
nm(I): print	names to user ID's
uids(V): map	new process
fork(II): create	nm(I): print namelist.
	octal dump of file
od(I):	od(I): octal dump of file
	off (format) text
roff(I): run	open file
close(II): close	open file
fstat(II): status of	open file
open(II):	open(II): open file
	(or print) files
cat(I): concatenate	output...a.out(V):
assembler and loader	owner of file
chown(II): change	owner of files
chown(I): change	paper tape format
bppt(V): binary	paper tape
dbppt(I): write binary	paper tape
lbppt(I): read binary	paper tape
ppt(IV): punched	paper tape

dli(VI):	load DEC binary
dpt(VI):	read DEC ASCII
as2(VII):	assembler's Fortran compiler
passwd(V):	paper tapes
mesg(I):	paper tapes
seek(II):	move read or write
tell(II):	find read or write
cal(VI):	pass 2
type(I):	passes...f1,f2,f3,f4(VII):
pr(I):	passwd(V): password file
cat(I):	password file
nm(I):	permit or deny messages
mesg(III):	pointer
ptime(III):	pointer
bproc(VII):	ppt(IV): punched paper tape
fork(II):	pr(I): print file with headings
init(VII):	print calendar
rele(II):	print file on IBM 2741
wait(II):	print file with headings
break(II):	print) files
exec(II):	print namelist
b(I):	print string on typewriter
for(I):	print time
ppt(IV):	procedure
quit(II):	process
lbppt(I):	process
dpt(VI):	processor
read(II):	process
seek(II):	program break
tell(II):	program file
boot(I):	program
rele(II):	program
strip(I):	ptime(III): print time
tty0,...,tty5(IV):	punched paper tape
umount(I):	putc(III): write character or word
rmdir(I):	quit(II): catch or inhibit quits
rm(I):	quits
unlink(II):	read binary paper tape
strip(I):	read DEC ASCII paper tapes
mv(I):	read file
rew(I):	read or write pointer
rf0(IV):	read or write pointer
rk0(IV):	read(II): read file
rkf(I):	reboot system
format	release processor
	rele(II): release processor
	relocation bits
	remote typewriter
	removable file system
	remove (delete) directory
	remove (delete) file
	remove (delete) file
	remove symbols, relocation bits
	rename file
	rew(I): rewind DECTape
	rewind DECTape
	RF disk file
	rf0(IV): RF disk file
	RK disk file
	RK disk
	rkd(I): dump disk to tape

```

rkf(I): format RK disk
rkl(I): load disk from tape
rk0(IV): RK disk file
rmdir(I): remove (delete) directory
rm(I): remove (delete) file
roff(I): run off (format) text
roff's suffix table
run off (format) text
sdate(I): adjust date and time .
seek(II): move read or write pointer
send mail to another user
set date modified of file
set mode of typewriter
set program break
set system time
set tab stops on typewriter
set user ID
setuid(II): set user ID
Shell
sh(I): command interpreter
simulator
sine, cosine
sin(III): sine, cosine
smdate(II): set date modified of file
sort a file
sort(VI): sort a file
space
standard assembly-language library
standard B library
standard Fortran library
start and finish
stat(I): get file status
stat(II): get file status
status of open file
status
status
stime(II): set system time
stops on typewriter
string on typewriter
strip(I): remove symbols, relocation bits
stty(II): set mode of typewriter
suffix table
su(I): become super-user
sum file
sum(I): sum file
super-user
supplied BASIC
switch(III): transfer depending on value
symbolic debugger
symbols, relocation bits
symbols
system format
system time
system
system
system...check(I):
check consistency of file

```

and logging off the
 mkfs(I): initialize file
 mount detachable file
 mount(II): mount file
 dismount removable file
 umount(II): dismount file
 . file
 who(I): who is on the
 tabs(VII): set
 suftab(VII): roff's suffix

 bppt(V): binary paper
 dbppt(I): write binary paper
 lbppt(I): read binary paper
 ppt(IV): punched paper
 rkd(I): dump disk to
 rkl(I): load disk from
 load DEC binary paper
 dpt(VI): read DEC ASCII paper

 tty(I): find name of
 exit(II):
 ed(I):
 roff(I): run off (format)
 ttt(VI): the game of
 tm(I): get
 date(I): get date and
 time(II): get
 ctime(III): convert

 ptime(III): print
 sdate(I): adjust date and
 stime(II): set system

 switch(III):
 catch illegal instruction
 cemt(II): catch EMT

 kbd(VII): map of

 gtty(II): get
 hup(I): hang up
 mesg(III): print string on
 stty(II): set mode of
 tabs(VII): set tab stops on
 tty(IV): console
 tty0,...,tty5(IV): remote

 un(I): find

system...login, logout(VII): logging on
 system
 system...mount(I):
 system
 system...umount(I):
 system
 system(V): file system format
 system
 tab stops on typewriter
 table
 tabs(VII): set tab stops on typewriter
 tape format
 tape
 tape
 tape
 tape
 tapes...dli(VI):
 tapes
 tap(I): manipulate DECTape
 tap0,...,tap7(IV): DECTape file
 tell(II): find read or write pointer
 terminal
 terminate execution
 text editor
 text
 tic-tac-toe
 time information
 time of day
 time of year
 time to ASCII
 time(II): get time of year
 time
 time
 time
 tm(I): get time information
 transfer depending on value
 trap...ilgins(II):
 traps
 ttt(VI): the game of tic-tac-toe
 TTY 37 keyboard
 tty(I): find name of terminal
 tty(IV): console typewriter
 tty0,...,tty5(IV): remote typewriter
 type(I): print file on IBM 2741
 typewriter mode
 typewriter
 typewriter
 typewriter
 typewriter
 typewriter
 typewriter
 uids(V): map names to user ID's
 umount(I): dismount removable file system
 umount(II): dismount file system
 undefined symbols

un(I): find undefined symbols
unlink(II): remove (delete) file
usage
user ID
user ID
user ID's
user information
user
user
utmp(V): logged-in user information
value...switch(III):
wait for process
wait(II): wait for process
wc(I): get (English) word count
who is on the system
who(I): who is on the system
with given name
with headings
word count
word
working directory
working directory
write binary paper tape
write character or word
write pointer
write pointer
write to another user
write(I): write to another user
year
2
2741
37 keyboard

du(I): find disk
getuid(II): get
setuid(II): set
uids(V): map names to
utmp(V): logged-in
mail(I): send mail to another
write(I): write to another
transfer depending on
wait(II):
who(I):
find(I): find file
pr(I): print file
wc(I): get (English)
putc(III): write character or
chdir(I): change
chdir(II): change
dbppt(I):
putc(III):
seek(II): move read or
tell(II): find read or
write(I):
time(II): get time of
as2(VII): assembler's pass
type(I): print file on IBM
kbd(VII): map of TTY

NAME ar -- archive

SYNOPSIS ar key afile name, ...

DESCRIPTION ar maintains groups of files combined into a single archive file. Its main use is to create and update library files as used by the loader. It can be used, though, for any similar purpose.

key is one character from the set drtux, optionally concatenated with y. afile is the archive file. The names are constituent files in the archive file. The meanings of the key characters are:

d means delete the named files from the archive file.

r means replace the named files in the archive file. If the archive file does not exist, r will create it. If the named files are not in the archive file, they are appended.

t prints a table of contents of the archive file. If no names are given, all files in the archive are tabbed. If names are given, only those files are tabbed.

u is similar to r except that only those files that have been modified are replaced. If no names are given, all files in the archive that have been modified will be replaced by the modified version.

x will extract the named files. If no names are given, all files in the archive are extracted. In neither case does x alter the archive file.

v means verbose. Under the verbose option, ar gives a file-by-file description of the making of a new archive file from the old archive and the constituent files. The following abbreviations are used:

<u>c</u>	copy
<u>a</u>	append
<u>d</u>	delete
<u>r</u>	replace
<u>x</u>	extract

FILES /tmp/vtma, vtmb ... temporary

SEE ALSO ld

DIAGNOSTICS "Bad usage", "afile -- not in archive format", "cannot open temp file", "name -- cannot open",

11/3/71

AR (I)

"name -- phase error", "name -- cannot create",
"no archive file", "cannot create archive file",
"name -- not found".

BUGS

Option l (table with more information) should be implemented.

There should be a way to specify the placement of a new file in an archive. Currently, it is placed at the end.

OWNER

ken, dmr

NAME **as** -- assembler

SYNOPSIS **as** **name**₁ ...

DESCRIPTION **as** assembles the concatenation of **name**₁, **as** is based on the DEC-provided assembler PAL-11R [references], although it was coded locally. Therefore, only the differences will be recorded.

Character changes are:

for	use
@	*
#	\$
;	/

In **as**, the character ";" is a logical new line; several operations may appear on one line if separated by ";". Several new expression operators have been provided:

\>	right shift (logical)
\<	left shift
*	multiplication
\vee	division
%	remainder (no longer means "register")
!	one's complement
[]	parentheses for grouping result has value of left, type of right

For example location 0 (relocatable) can be written "0^".; another way to denote register 2 is "2^r0".

All of the preceding operators are binary; if a left operand is missing, it is taken to be 0. The "!" operator adds its left operand to the one's complement of its right operand.

There is a conditional assembly operation code different from that of PAL-11R (whose conditionals are not provided):

```
.if expression
...
.endif
```

If the expression evaluates to non-zero, the section of code between the ".if" and the ".endif" is assembled; otherwise it is ignored. ".if"s may be nested.

Temporary labels like those introduced by Knuth [reference] may be employed. A temporary label is defined as follows:

n:

where n is a digit 0 ... 9. Symbols of the form "nf" refer to the first label n: following the use of the symbol; those of the form "nb" refer to the last n:. The same n may be used many times. Labels of this form are less taxing both on the imagination of the programmer and on the symbol table space of the assembler.

The PAL-11R opcodes ".eot" and ".end" are redundant and are omitted.

The symbols

```
r0 ... r5
sp
pc
ac
mq
div
mul
lsh
ash
nor
csw
..
..
```

are predefined with appropriate values. The symbol "csw" refers to the console switches. .. is the relocation constant and is added to each relocatable symbol; normally it is 40000(8); it may be changed to assemble a section of code at a location different from that in which it will be executed.

It is illegal to assign a value to "." less than its current value.

The new opcode "sys" is used to specify system calls. Names for system calls are predefined. See the section on system calls for their names.

Strings of characters may be assembled in a way more convenient than PAL-11's ".ascii" operation (which is, therefore, omitted). Strings are included between the string quotes "<" and ">":

<here is a string>

Escape sequences exist to enter non graphic and other difficult characters. These sequences are also effective in single and double character constants introduced by single ('') and double ("") quotes respectively.

```

use  for
\n  newline (012)
\0  NULL (000)
\>  >
\t  TAB (011)
\\  \

```

The binary output of the assembler is placed on the file "a.out" in the current directory. a.out also contains the symbol table from the assembly and relocation bits. The output of the assembler is executable immediately if the assembly was error-free and if there were no unresolved external references. The link editor ld may be used to combine several assembly outputs and resolve global symbols.

The multiple location counter feature of PAL11R is not supported.

The assembler does not produce a listing of the source program. This is not a serious drawback; the debugger db discussed below is sufficiently powerful to render a printed octal translation of the source unnecessary.

FILES	/etc/as2	pass 2 of the assembler
	a.tmp1	temporary
	a.tmp2	temporary
	a.tmp3	temporary
	a.out	object

SEE ALSO ld, nm, sh, un, db, a.out (format of output)

DIAGNOSTICS When an input file cannot be read, its name followed by a question mark is typed and assembly ceases.

When syntactic or semantic errors occur, a single-character diagnostic is typed out together with the line number and the file name in which it occurred. Errors in pass 1 cause cancellation of pass 2. The possible errors are:

)	parentheses error
]	parentheses error
*	Indirection (" <u>*</u> ") used illegally
A	error in <u>Address</u>
B	Branch instruction has too remote an address
E	error in <u>Expression</u>
F	error in local (" <u>F</u> " or "b") type symbol
G	Garbage (unknown) character
M	Multiply defined symbol as label
O	<u>Odd</u> -- word quantity assembled at odd

11/3/71

AS (I)

address
P Phase error-- " . " different in pass 2 from
pass 1 value
R Relocation error
U Undefined symbol
X syntax error

- BUGS Symbol table overflow is not checked.

OWNER dmr

NAME B -- language

SYNOPSIS sh rc /usr/b/rc name

DESCRIPTION B is a language suitable for system programming.
It is described in a separate publication B
reference manual.

The canned shell sequence in /usr/b/rc will com-
pile the program name.b into the executable file
a.out. It involves running the B compiler, the B
assembler, the assembler and the link editor.
The process leaves the files name.i and name.s in
the current directory.

FILES name.b, name.i, name.s.

SEE ALSO /etc/bc, /etc/ba, /etc/brt1, /etc/brt2,
 /etc/bilib, /etc/libb.a, B reference manual.

DIAGNOSTICS see B reference manual

BUGS There should be a B command.

OWNER ken, dmr

NAME bas -- basic

SYNOPSIS bas [file]

DESCRIPTION bas is a dialect of basic. If a file argument is provided, the file is used for input before the console is read.

bas accepts lines of the form:

```
statement
integer statement
```

Integer numbered statements (known as internal statements) are stored for later execution. They are stored in sorted ascending order. Non-numbered statements are immediately executed. The result of an immediate expression statement (that does not have '=' as its highest operator) is printed.

Statements have the following syntax: (expr is short for expression)

expr

The expression is executed for its side effects (assignment or function call) or for printing as described above.

done

Return to system level.

draw expr expr expr

draw is used to draw on a 611-type storage scope through a TSP-1 plotter interface. The coordinates of the scope face are zero to one in both the x and y directions. (Zero,zero being the lower left corner.) The expressions are evaluated and designated X, Y, and Z. A line is drawn from the previous X, Y to the new X, Y. If Z is non-zero, the line is visible, otherwise the line is invisible.

for name = expr expr statement

for name = expr expr

...

next

The for statement repetatively executes a statement (first form) or a group of statements (second form) under control of a named variable. The variable takes on the value of the first expression, then is incremented by one on each loop, not to exceed the value of the second expression.

goto expr

The expression is evaluated, truncated to an integer and execution goes to the corresponding integer numbered statement. If executed from immediate mode, the internal statements are compiled first.

if expr statement

The statement is executed if the expression evaluates to non-zero.

list [expr [expr]]

list is used to print out the stored internal statements. If no arguments are given, all internal statements are printed. If one argument is given, only that internal statement is listed. If two arguments are given, all internal statements inclusively between the arguments are printed.

print expr

The expression is evaluated and printed.

return expr

The expression is evaluated and the result is passed back as the value of a function call.

run

The internal statements are compiled. The symbol table is re-initialized. The random number generator is re-set. Control is passed to the lowest numbered internal statement.

Expressions have the following syntax:

name

A name is used to specify a variable. Names are composed of a letter ('a' - 'z') followed by letters and digits. The first four characters of a name are significant.

number

A number is used to represent a constant value. A number is composed of digits, at most one decimal point ('.') and possibly a scale factor of the form e digits or e- digits.

{ expr }

Parentheses are used to alter normal order of evaluation.

expr op expr

Common functions of two arguments are

abbreviated by the two arguments separated by an operator denoting the function. A complete list of operators is given below.

`expr [expr [, expr ...]]`

Functions of an arbitrary number of arguments can be called by an expression followed by the arguments in parentheses separated by commas. The expression evaluates to the line number of the entry of the function in the internally stored statements. This causes the internal statements to be compiled. If the expression evaluates negative, an builtin function is called. The list of builtin functions appears below.

`name [expr [, expr ...]]`

Arrays are not yet implemented.

The following is the list of operators:

=

= is the assignment operator. The left operand must be a name or an array element. The result is the right operand. Assignment binds right to left, all other operators bind left to right.

& |

& (logical and) has result zero if either of its arguments are zero. It has result one if both its arguments are non-zero. | (logical or) has result zero if both of its arguments are zero. It has result one if either of its arguments are non-zero.

< <= > >= == <>

The relational operators (< less than, <= less than or equal, > greater than, >= greater than or equal, == equal to, <> not equal to) return one if their arguments are in the specified relation. They return zero otherwise. Relational operators at the same level extend as follows: a>b>c is the same as a>b&b>c.

+ -

Add and subtract.

* /

Multiply and divide.

^

Exponeniation.

The following is a list of builtin functions:

arg
Arg(i) is the value of the ith actual parameter on the current level of function call.

exp
Exp(x) is the exponential function of x.

log
Log(x) is the logarithm bas e of x.

sin
Sin(x) is the sine of x (radians).

cos
Cos(x) is the cosine of x (radians).

atn
Atn(x) is the arctangent of x. (Not implemented.)

rnd
Rnd() is a uniformly distributed random number between zero and one.

expr
Expr() is the only form of program input. A line is read from the input and evaluated as an expression. The resultant value is returned.

int
Int(x) returns x truncated to an integer.

FILES /tmp/btma, btmb ... temporary

SEE ALSO --

DIAGNOSTICS Syntax errors cause the incorrect line to be typed with an underscore where the parse failed. All other diagnostics are self explanatory.

BUGS Arrays [] are not yet implemented. In general, program sizes, recursion, etc are not checked, and cause trouble.

OWNER ken

NAME **bcd** — binary coded decimal conversion

SYNOPSIS **bcd** [string]

DESCRIPTION **bcd** will convert a string into GECOS card code.
If no argument string is provided, **bcd** will read
a line and convert it.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER dmr

NAME boot -- reboot system

SYNOPSIS /etc/boot

DESCRIPTION boot logically a command, and is kept in /etc only to lessen the probability of its being invoked by accident or from curiosity. It reboots the system by jumping to the read-only memory, which contains a disk boot program.

FILES --

SEE ALSO boot procedure

DIAGNOSTICS --

BUGS Should obviously not be executable by the general user. Also, it should reboot in a more direct manner. The mechanism invoked by jumping to the ROM loader is sensitive to the contents of the console switches, which makes the whole procedure even more dangerous.
Rather than jumping to the ROM, boot should simulate the ROM action with 173700 in the switches. In this manner, It may be used when the switches are not set, and even in installation without a ROM.

OWNER ken

NAME **cat** -- concatenate and print

SYNOPSIS **cat** file, ...

DESCRIPTION **cat** reads each file in sequence and writes it on the standard output stream. Thus:

cat file

is about the easiest way to print a file. Also:

cat file1 file2 >file3

is about the easiest way to concatenate files.

If no input file is given **cat** reads from the standard input file.

FILES --

SEE ALSO pr, cp

DIAGNOSTICS none; if a file cannot be found it is ignored.

BUGS --

OWNER ken, dmr

NAME **chdir** -- change working directory

SYNOPSIS **chdir** **directory**

DESCRIPTION **directory** becomes the new working directory.
Because a new process is created to execute each command, **chdir** would be ineffective if it were written as a normal command. It is therefore recognized and executed by the Shell.

FILES --

SEE ALSO **sh**

DIAGNOSTICS ?

BUGS --

OWNER ken, dmr

NAME **check** -- file system consistency check

SYNOPSIS . **check** [filesystem [blockno, ...]]

DESCRIPTION **check** will examine a file system, build a bit map of used blocks, and compare this bit map against the bit map maintained on the file system. If the file system is not specified, a check of both /dev/rf0 and /dev/rk0 is performed. Output includes the number of files on the file system, the number of these that are 'large', the number of used blocks, and the number of free blocks.

FILES /dev/rf0, /dev/rk0

SEE ALSO find

DIAGNOSTICS Diagnostics are produced for blocks missing, duplicated, and bad block addresses. Diagnostics are also produced for block numbers passed as parameters. In each case, the block number, i-number, and block class (i = inode, x indirect, f free) is printed.

BUGS The checking process is two pass in nature. If checking is done on an active file system, extraneous diagnostics may occur.
The swap space on the RF file system is not accounted for and will therefore show up as 'missing'.

OWNER ken, dmr

NAME **chmod** -- change mode

SYNOPSIS **chmod** octal file, ...

DESCRIPTION The octal mode replaces the mode of each of the files. The mode is constructed from the OR of the following modes:

01 write for non-owner
02 read for non-owner
04 write for owner
10 read for owner
20 executable
40 set-UID

Only the owner of a file may change its mode.

FILES --

SEE ALSO stat, ls

DIAGNOSTICS ?

BUGS --

OWNER ken, dmr

NAME **chown** -- change owner

SYNOPSIS **chown** owner file, ...

DESCRIPTION owner becomes the new owner of the files. The owner may be either a decimal UID or a name found in /etc/uids.

Only the owner of a file is allowed to change the owner. It is illegal to change the owner of a file with the set-user-ID mode.

FILES /etc/uids

SEE ALSO **stat**

DIAGNOSTICS ?

BUGS --

OWNER ken, dmr

11/3/71

CMP (I)

NAME **cmp** -- compare two files

SYNOPSIS **cmp** file₁ file₂

DESCRIPTION The two files are compared for identical contents. Discrepancies are noted by giving the offset and the differing words.

FILES --

SEE ALSO --

DIAGNOSTICS Messages are given for inability to open either argument, premature EOF on either argument, and incorrect usage.

BUGS If the two files differ in length by one byte, the extra byte does not enter into the comparison.

OWNER dmr

NAME **cp** -- copy

SYNOPSIS **cp** file₁₁ file₁₂ file₂₁ file₂₂ ...

DESCRIPTION Files are taken in pairs; the first is opened for reading, the second created mode 17. Then the first is copied into the second.

FILES --

SEE ALSO cat, pr

DIAGNOSTICS Error returns are checked at every system call, and appropriate diagnostics are produced.

BUGS The second file should be created in the mode of the first.
A directory convention as used in **mv** should be adopted to **cp**.

OWNER ken, dmr

11/3/71

DATE (I)

NAME date -- print the date
SYNOPSIS date
DESCRIPTION The current date is printed to the second.
FILES --
SEE ALSO sdate
DIAGNOSTICS --
BUGS --
OWNER dmr

NAME db -- debug

SYNOPSIS db [core [namelist]]

DESCRIPTION Unlike many debugging packages (including DEC's ODT, on which db is loosely based) db is not loaded as part of the core image which it is used to examine; instead it examines files. Typically, the file will be either a core image produced after a fault or the binary output of the assembler. Core is the file being debugged; if omitted core is assumed. namelist is a file containing a symbol table. If it is omitted, a.out is the default. If no appropriate name list file can be found, db can still be used but some of its symbolic facilities become unavailable.

The format for most db requests is an address followed by a one character command.

Addresses are expressions built up as follows:

1. A name has the value assigned to it when the input file was assembled. It may be relocatable or not depending on the use of the name during the assembly.
2. An octal number is an absolute quantity with the appropriate value.
3. An octal number immediately followed by "r" is a relocatable quantity with the appropriate value.
4. The symbol "." indicates the current pointer of db. The current pointer is set by many db requests.
5. Expressions separated by "+" or " " (blank) are expressions with value equal to the sum of the components. At most one of the components may be relocatable.
6. Expressions separated by "-" form an expression with value equal to the difference to the components. If the right component is relocatable, the left component must be relocatable.
7. Expressions are evaluated left to right.

Names for registers are built in:

r0 ... r5
sp
pc

ac
mq

These may be examined. Their values are deduced from the contents of the stack in a core image file. They are meaningless in a file that is not a core image.

If no address is given for a command, the current address (also specified by ".") is assumed. In general, "." points to the last word or byte printed by db.

There are db commands for examining locations interpreted as octal numbers, machine instructions, ASCII characters, and addresses. For numbers and characters, either bytes or words may be examined. The following commands are used to examine the specified file.

- / The addressed word is printed in octal.
 - \ The addressed byte is printed in octal.
 - " The addressed word is printed as two ASCII characters.
 - ' The addressed byte is printed as an ASCII character.
 - * The addressed word is multiplied by 2, then printed in octal (used with B programs, whose addresses are word addresses).
 - ? The addressed word is interpreted as a machine instruction and a symbolic form of the instruction, including symbolic addresses, is printed. Usually, the result will appear exactly as it was written in the source program.
 - & The addressed word is interpreted as a symbolic address and is printed as the name of the symbol whose value is closest to the addressed word, possibly followed by a signed offset.
- <nl> (i. e., the character "new line") This command advances the current location counter "." and prints the resulting location in the mode last specified by one of the above requests.
- ^ This character decrements "." and prints the resulting location in the mode last selected one of the above requests. It is

a converse to <nl>.

It is illegal for the word-oriented commands to have odd addresses. The incrementing and decrementing of ". " done by the <nl> and requests is by one or two depending on whether the last command was word or byte oriented.

The address portion of any of the above commands may be followed by a comma and then by an expression. In this case that number of sequential words or bytes specified by the expression is printed. ". " is advanced so that it points at the last thing printed.

There are two commands to interpret the value of expressions.

- = When preceded by an expression, the value of the expression is typed in octal. When not preceded by an expression, the value of ". " is indicated. This command does not change the value of ". ".
- : An attempt is made to print the given expression as a symbolic address. If the expression is relocatable, that symbol is found whose value is nearest that of the expression, and the symbol is typed, followed by a sign and the appropriate offset. If the value of the expression is absolute, a symbol with exactly the indicated value is sought and printed if found; if no matching symbol is discovered, the octal value of the expression is given.

The following command may be used to patch the file being debugged.

- ! This command must be preceded by an expression. The value of the expression is stored at the location addressed by the current value of ". ". The opcodes do not appear in the symbol table, so the user must assemble them by hand.

The following command is used after a fault has caused a core image file to be produced.

- \$ causes the contents of the general registers and several other registers to be printed both in octal and symbolic format. The values are as they were at the time of the fault.

The only way to exit from db is to generate an end of file on the typewriter (EOT character).

FILES

--

SEE ALSO

as; core for format of core image.

DIAGNOSTICS

"File not found" if the first argument cannot be read; otherwise "?".

BUGS

Really, db should know about relocation bits, floating point operations, and PDP11/45 instructions.

OWNER

dmr

NAME dbppt -- dump binary paper tape

SYNOPSIS dbppt name [output]

DESCRIPTION dbppt produces binary paper tape in UNIX standard format, which includes checksums and a zero-suppression feature. File name is dumped; if the output argument is not given, output goes to /dev/ppt.

FILES /dev/ppt

SEE ALSO lbppt to reload the tapes. bppt for binary paper tape format.

DIAGNOSTICS ?

BUGS --

OWNER ken

NAME	dc -- desk calculator
SYNOPSIS	<u>dc</u>
DESCRIPTION	<u>dc</u> is an arbitrary precision integer arithmetic package. The overall structure of dc is a stacking (reverse Polish) calculator. The following constructions are recognized by the calculator:

number
The value of the number is pushed on the stack. If the number starts with a zero, it is taken to be octal, otherwise it is decimal.

t = * / %
The top two values on the stack are added (+), subtracted (-), multiplied (*), divided (/), or remaindered (%). The two entries are popped off of the stack, the result is pushed on the stack in their place.

sx
The top of the stack is popped and stored into a register named x, where x may be any character.

lx
The value in register x is pushed on the stack. The register x is not altered.

d
The top value on the stack is pushed on the stack. Thus the top value is duplicated.

p
The top value on the stack is printed in decimal. The top value remains unchanged.

f
All values on the stack are popped off and printed in decimal.

r
All values on the stack are popped.

q
exit.

h
print brief synopsis of commands to dc.

new-line
space
ignored.

An example to calculate the monthly, weekly and

11/3/71

DC (I)

hourly rates for a \$10,000/year salary.

10000
100* (now in cents)
dsa (non-destructive store)
12/ (pennies per month)
la52/ (pennies per week)
d10* (deci-pennies per week)
375/ (pennies per hour)
f (print all results)
(3) 512
(2) 19230
(1) 83333

FILES --

SEE ALSO --

DIAGNOSTICS ? (x) for unrecognized character x.

BUGS % doesn't work correctly.

OWNER ken

NAME df -- disk free

SYNOPSIS df [filesystem]

DESCRIPTION df prints out the number of free blocks available on a file system. If the file system is unspecified, the free space on /dev/rf0 and /dev/rk0 is printed.

FILES /dev/rf0, /dev/rk0

SEE ALSO check

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME dsw -- delete interactively

SYNOPSIS dsw [directory]

DESCRIPTION For each file in the given directory ("." if not specified) dsw types its name. If "y" is typed, the file is deleted; if "x", dsw exits; if anything else, the file is not removed.

FILES --

SEE ALSO rm

DIAGNOSTICS "?"

BUGS The name "dsw" is a carryover from the ancient past. Its etymology is amusing but the name is nonetheless ill-advised.

OWNER dmr, ken

NAME dtf -- DECtape format

SYNOPSIS /etc/dtf

DESCRIPTION dtf will write timing tracks, mark tracks and block numbers on a virgin DECtape. The format is DEC standard of 578 blocks of 256 words each. The end zones are a little longer than standard DEC.

Before use, the tape to be formatted should be mounted on drive 0. The 'wall' and 'wtm' switches should be enabled. After the tape is formatted, the switches should be disabled to prevent damage to subsequent tapes due to a controller logic error.

FILES --

SEE ALSO sdate

DIAGNOSTICS "?" is typed for any error detected.

BUGS This program does physical I/O on drive 0. The processor priority is set very high due to very stringent real time requirements. This means that all time sharing activities are suspended during the formatting (about 1.5 minutes) The real time clock will also be slow.

OWNER ken

NAME du -- summarize disk usage

SYNOPSIS du [-s] [-a] [name ...]

DESCRIPTION du gives the number of blocks contained in all files and (recursively) directories within each specified directory or file name. If name is missing, . is used.

The optional argument -s causes only the grand total to be given. The optional argument -a causes an entry to be generated for each file. Absence of either causes an entry to be generated for each directory only.

A file which has two links to it is only counted once.

FILES /

SEE ALSO --

DIAGNOSTICS --

BUGS Files at the top level (not under -a option) are not listed.

Removable file systems do not work correctly since i-numbers may be repeated while the corresponding files are distinct. Du should maintain an i-number list per root directory encountered.

OWNER dmr

NAME ed -- editor

SYNOPSIS ed [name]

DESCRIPTION ed is the standard text editor. ed is based on QED [reference] but is fully if succinctly described here. Differences between ed and QED are also noted to simplify the transition to the less powerful editor.

If the optional argument is given, ed simulates an e command on the named file; that is to say, the file is read into ed's buffer so that it can be edited.

ed operates on a copy of any file it is editing; changes made in the copy have no effect on the file until an explicit write (w) command is given. The copy of the text being edited resides in a temporary file called the buffer. There is only one buffer.

Commands to ed have a simple and regular structure: zero or more addresses followed by a single character command, possibly followed by parameters to the command. These addresses specify one or more lines in the buffer. Every command which requires addresses has default addresses, so that the addresses can often be omitted.

In general only one command may appear on a line. Certain commands allow the input of text. This text is placed in the appropriate place in the buffer. While ed is accepting text, it is said to be in input mode. In this mode, no commands are recognized; all input is merely collected. Input mode is left by typing a period (.) alone at the beginning of a line.

ed supports a limited form of regular expression notation. A regular expression is an expression which specifies a set of strings of characters. A member of this set of strings is said to be matched by the regular expression. The regular expressions allowed by ed are constructed as follows:

1. An ordinary character (not one of those discussed below) is a regular expression and matches that character.
2. A circumflex (^) at the beginning of a regular expression matches the null character at the beginning of a line.

3. A currency symbol (\$) at the end of a regular expression matches the null character at the end of a line.
4. A period (.) matches any character but a new-line character.
5. A regular expression followed by an asterisk (*) matches any number of adjacent occurrences (including zero) of the regular expression it follows.
6. A string of characters enclosed in square brackets ([]) matches any character in the string but no others. If, however, the first character of the string is a circumflex (^) the regular expression matches any character but new-line and the characters in the string.
7. The concatenation of regular expressions is a regular expression which matches the concatenation of the strings matched by the components of the regular expression.
8. The null regular expression standing alone is equivalent to the last regular expression encountered.

Regular expressions are used in addresses to specify lines and in one command (s, see below) to specify a portion of a line which is to be replaced.

If it is desired to use one of the regular expression metacharacters as an ordinary character, that character may be preceded by "\". This also applies to the character bounding the regular expression (often "/") and to "\\" itself.

Addresses are constructed as follows. To understand addressing in ed it is necessary to know that at any time there is a current line. Generally speaking, the current line is the last line affected by a command; however, the exact effect on the current line by each command is discussed under the description of the command.

1. The character "." addresses the current line.
2. The character "\$" addresses the last line of the buffer.
3. A decimal number n addresses the nth line of the buffer.

4. A regular expression enclosed in slashes "/" addresses the first line found by searching toward the end of the buffer and stopping at the first line containing a string matching the regular expression. If necessary the search wraps around to the beginning of the buffer.
5. A regular expression enclosed in queries "?" addresses the first line found by searching toward the beginning of the buffer and stopping at the first line found containing a string matching the regular expression. If necessary the search wraps around to the end of the buffer.
6. An address followed by a plus sign "+" or a minus sign "-" followed by a decimal number specifies that address plus (resp. minus) the indicated number of lines. The plus sign may be omitted.

Commands may require zero, one, or two addresses. Commands which require no addresses regard the presence of an address as an error. Commands which require the presence of one address all assume a default address (often ".") but if given more than one address ignore any extras and use the last given. Commands which require two addresses have defaults in the case of zero or one address but use the last two if more than two are given.

Addresses are separated from each other typically by a comma (,),. They may also be separated by a semicolon (;). In this case the current line "." is set to the previous address before the next address is interpreted. This feature is used to control the starting line for forward and backward searches ("/", "?").

In the following list of ed commands, the default addresses are shown in parentheses. The parentheses are not part of the address, but are used to show that the given addresses are the default.

As mentioned, it is generally illegal for more than one command to appear on a line. However, any command may be suffixed by "p" (for "print"). In that case, the current line is printed after the command is complete.

In any two-address command, it is illegal for the

first address to lie after the second address.

(.)a
<text>

- The append command reads the given text and appends it after the addressed line. " ." is left on the last line input, if there were any, otherwise at the addressed line. Address "0" is legal for this command; text is placed at the beginning of the buffer. (NOTE: the default address differs from that of QED.)

(...)c
<text>

- The change command deletes the addressed lines, then accepts input text which replaces these lines. " ." is left at the last line input; if there were none, it is left at the first line not changed.

(...)d

The delete command deletes the addressed lines from the buffer. " ." is left at the first line not deleted.

e filename

The edit command causes the entire contents of the buffer to be deleted, and then the named file to be read in. " ." is set to the last line of the buffer. The number of characters read is typed.

(1,\$)g/regular expression/command

In the global command, the first step is to mark every line which matches the given regular expression. Then for every such line, the given command is executed with " ." set to that line. The repeated command cannot be a, g, i, or c.

(.)i
<text>

- This command inserts the given text before the addressed line. " ." is left at the last line input; if there were none, at the addressed line. This command differs from the a command only in the placement of the text.

(...)l

The list command prints the addressed lines in an unambiguous way. Non-printing

characters are over-struck as follows:

char	prints
bs	\backslash
tab	\rightarrow
ret	\leftarrow
SI	\pm
SO	Θ

All characters preceded by a prefix (ESC) character are printed over-struck with without the prefix. Long lines are folded with the sequence \newline.

(.,.)p

The print command prints the addressed lines. $"."$ is left at the last line printed.

q

The quit command causes ed to exit. No automatic write of a file is done.

(\$)r filename

The read command reads in the given file after the addressed line. If no file name is given, the file last mentioned in e, r, or w commands is read. Address $"0"$ is legal for r and causes the file to be read at the beginning of the buffer. If the read is successful, the number of characters read is typed. $"."$ is left at the last line of the file.

(.,.)s/regular expression/replacement/

The substitute command searches each addressed line for an occurrence of the specified regular expression. On each line in which a match is found, the first (and only first, compare QED) matched string is replaced by the replacement specified. It is an error for the substitution to fail on all addressed lines. Any character other than space or new-line may be used instead of $/$ to delimit the regular expression and the replacement. $"."$ is left at the last line substituted.

The ampersand $"&"$ appearing in the replacement is replaced by the regular expression that was matched. The special meaning of $"&"$ in this context may be suppressed by preceding it by $"\\"$.

(1,\$)w filename

The write command writes the addressed lines onto the given file. If no file name is given, the file last named in e, r, or w

commands is written. " ." is unchanged. If the command is successful, the number of characters written is typed.

`(\$)=`

The line number of the addressed line is typed. " ." is unchanged by this command.

`!UNIX command`

The remainder of the line after the "!" is sent to UNIX to be interpreted as a command. " ." is unchanged.

`<newline>`

A blank line alone is equivalent to ".+1p"; it is useful for stepping through text.

Ed can edit at most 1500 lines and the maximum size of a line is 256 characters. The differences between ed and QED are:

1. There is no "\f" character; input mode is left by typing " ." alone on a line.
2. There is only one buffer and hence no "\b" stream directive.
3. The commands are limited to:

a c d e g i l p q r s w = !

where e is new.

4. The only special characters in regular expressions are:

* ^ \$ [.

which have the usual meanings. However, " ^ " and " \$" are only effective if they are the first or last character respectively of the regular expression. Otherwise suppression of special meaning is done by preceding the character by "\", which is not otherwise special.

5. In the substitute command, only the left-most occurrence of the matched regular expression is substituted.
7. The a command has a different default address.

FILES

/tmp/etma, etmb, ... temporary
 /etc/msh is used to implement the "!" command.

11/3/71

ED (I)

SEE ALSO

--

DIAGNOSTICS

"?" for any error

BUGS

ed is used as the shell for the editing system.
It has the editing system UID built in and if
invoked under this UID will give slightly dif-
ferent responses. This is a little kludgy.

OWNER

ken

11/3/71

FIND (I)

NAME find -- find file with given name

SYNOPSIS find name or number ...

DESCRIPTION find searches the entire file system hierarchy
and gives the path names of all files with the
specified names or (decimal) i-numbers.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER dmr

NAME for -- fortran

SYNOPSIS for file

DESCRIPTION for is a nearly complete fortran compiler. file is the name of a fortran source program to be compiled. The following is a list of differences between for and ANSI standard fortran:

1. arbitrary combination of types are allowed in expressions. Not all combinations are expected to be supported in runtime. All of the normal conversions involving integer, real and double precision are allowed.

FILES f.tmp1, 2 3 temporary
/etc/f1, 2 3 4 passes
/etc/xx runtime

SEE ALSO --

DIAGNOSTICS Diagnostics are given by number. If the source code is available, it is printed with an underline at the current character pointer. A listing of error numbers is available.

BUGS The following is a list of those features not yet implemented:

functions
arithmetic statement functions
data statements
complex constants
hollerith constants
continuation cards

OWNER dmr, ken

NAME	form -- form letter generator
SYNOPSIS	. <u>form</u> proto arg ₁ ...
DESCRIPTION	<u>form</u> generates a form letter from a prototype letter, an associative memory, arguments and in a special case, the current date. If <u>form</u> is invoked with the argument <u>x</u> , the following files come into play: x.f prototype input x.r form letter output x.am associative memory form.am associative memory if x.am not found.
	Basically, <u>form</u> is a copy process from the file x.f to the file x.r. If an element of the form \n (where n is a digit from 1 to 9) is encountered, The nth argument is inserted in its place, and that argument is then rescanned. If \0 is encountered, the current date is inserted. If the desired argument has not been given, a message of the form "\n:" is typed. The response typed in then is used for that argument.
	If an element of the form [name] is encountered, the name is looked up in the associative memory. If it is found, the contents of the memory under this name replaces the original element (again rescanned.) If the name is not found, a message of the form "name:" is typed. The response typed in is used for that element. If the associative memory is writable, the response is entered in the memory under the name. Thus the next search for that name will succeed without interaction.
	In both of the above cases, the response is typed in by entering arbitrary text terminated by two new lines. Only the first of the two new lines is passed with the text. The process is instantly terminated if an end of file is encountered anywhere except in the associative memory.
FILES	x.f input file x.r output file x.am associative memory form.am associative memory
SEE ALSO	type
DIAGNOSTICS	"setup error" when the appropriate files cannot be located or created.
BUGS	"setup" is misspelled.

11/3/71

FORM (I)

OWNER

rhm, ken

11/3/71

HUP (I)

NAME hup -- hang up typewriter

SYNOPSIS hup

DESCRIPTION hup hangs up the phone on the typewriter which uses it.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS should not be used; sometimes causes the typewriter channel to be lost.

OWNER dmr, ken

NAME **lbppt** -- load binary paper tapes

SYNOPSIS **lbppt** output [input]

DESCRIPTION **lbppt** loads a paper tape in standard UNIX binary paper tape format. It is used to bring files to a UNIX installation. Currently there is a GECOS program to prepare a GECOS file in binary paper tape format.

If the input file is specified, the character stream from that input is expected to be in UNIX binary paper tape format. If it is not present, /dev/ppt is assumed. The input stream is interpreted, checksummed, and copied to the output file.

FILES /dev/ppt

SEE ALSO dbppt, bppt format

DIAGNOSTICS "checksum"; "usage: "; "read error".

BUGS --

OWNER ken

NAME `ld` -- link editor

SYNOPSIS `ld [-usaol] name,]`

DESCRIPTION `ld` combines several object programs into one; resolves external references; and searches libraries. In the simplest case the names of several object programs are given, and `ld` combines them, producing an object module which can be either executed or become the input for a further `ld` run.

The argument routines are concatenated in the order specified. The entry point of the output is the beginning of the first routine.

If any argument is a library, it is searched, and only those routines defining an unresolved external reference are loaded. If any routine loaded from a library refers to an undefined symbol which does not become defined by the end of the library, the library is searched again. Thus the order of libraries primarily affects the efficiency of loading, not what routines get loaded.

`ld` understands several flag arguments which are written preceded by a "`-`":

- s "squash" the output, that is, remove the symbol table and relocation bits to save space (but impair the usefulness of the debugger). This information can also be removed by strip.
- u take the following argument as a symbol and enter it as undefined in the symbol table. This is useful for loading wholly from a library, since initially the symbol table is empty and an unresolved reference is needed to force the loading of the first routine.
- o set the origin of the load to the octal number which is given as the next argument. This option affects only the definition of relocatable external symbols. See DMR before using.
- l This option is an abbreviation for a library name. "`-l`" alone stands for "/etc/liba.a", which is the standard system library for assembly language programs. "`-lx`" stands for "/etc/libx.a" where x is any character. There are libraries for Fortran (`x="f"`) and B (`x="b"`).

-a means "absolute" (load at origin absolute 0) but it doesn't work.

The output of ld is left on a.out. This file is executable only if no errors occurred during the load.

FILES

/etc/libx.a, for various x;
/etc/ltma, ltmb, ... (temporary)
a.out (output file)

SEE ALSO

as, strip, ar (maintains libraries)

DIAGNOSTICS

"can't create temp file"-- unwritable directory or someone else is using ld in the same directory.

"can't open temp file"-- maybe someone has deleted it out from under you.

"file not found"-- bad argument

"bad format"-- bad argument

"relocation error"-- bad argument (relocation bits corrupted)

"bad relocation"-- user error: a relocatable reference to an external symbol that turns out to be absolute.

"multiply defined"-- same symbol defined twice in same load

"un"-- stands for "undefined symbol"

"symbol not found"-- loader bug

BUGS

Option "-a" doesn't work at all; option "-o" doesn't work right.

OWNER

dmr

NAME ln -- make a link

SYNOPSIS ln name₁ [name₂]

DESCRIPTION ln creates a link to an existing file name₁. If name₂ is given, the link has that name; otherwise it is placed in the current directory and its name is the last component of name₁.
It is forbidden to link to a directory or to link across file systems.

FILES --

SEE ALSO rm, to unlink

DIAGNOSTICS "?"

BUGS There is nothing particularly wrong with ln, but links don't work right with respect to the backup system: one copy is backed up for each link, and (more serious) in case of a file system reload both copies are restored and the information that a link was involved is lost.

OWNER ken, dmr

NAME ls -- list contents of directory

SYNOPSIS ls [-ltasd] name, ...

DESCRIPTION ls lists the contents of one or more directories under control of several options:

 l list in long format, giving i-number, mode, owner, size in bytes, and time of last modification for each file. (see stat for format of the mode)

 t sort by time modified (latest first) instead of by name, as is normal

 a list all entries; usually those beginning with "." are suppressed

 s give size in blocks for each entry

 d if argument is a directory, list only its name, not its contents (mostly used with "-l" to get status on directory)

If no argument is given, "." is listed. If an argument is not a directory, its name is given.

FILES /etc/uids to get user ID's for ls -l

SEE ALSO stat

DIAGNOSTICS "name nonexistent"; "name unreadable"; "name unstatale."

BUGS In ls -l, when a user cannot be found in /etc/uids, the user number printed instead of a name is incorrect. It is correct in stat.

OWNER dmr, ken

NAME mail -- send mail to another user

SYNOPSIS mail [letter person ...]

DESCRIPTION mail without an argument searches for a file called mailbox, prints it if present, and asks if it should be saved. If the answer is "y", the mail is renamed mail, otherwise it is deleted. The answer to the above question may be supplied in the letter argument.

When followed by the names of a letter and one or more people, the letter is appended to each person's mailbox. Each letter is preceded by the sender's name and a postmark.

A person is either the name of an entry in the directory /usr, in which case the mail is sent to /usr/person/mailbox, or the path name of a directory, in which case mailbox in that directory is used.

When a user logs in he is informed of the presence of mail.

FILES /etc/uids to map the sender's numerical user ID to name; mail and mailbox in various directories.

SEE ALSO init

DIAGNOSTICS "Who are you?" if the user cannot be identified for some reason (a bug). "Cannot send to user" if mailbox cannot be opened.

BUGS --

OWNER ken

NAME mesg -- permit or deny messages

SYNOPSIS mesq [n][y]

DESCRIPTION mesq n forbids messages via write by revoking non-user write permission on the user's typewriter. mesq y reinstates permission. mesq with no argument reverses the current permission. In all cases the previous state is reported.

FILES /dev/ttyn

SEE ALSO write

DIAGNOSTICS "?" if the standard input file is not a typewriter

BUGS --

OWNER dmr, ken

11/3/71

MKDIR (I)

NAME mkdir -- make a directory

SYNOPSIS mkdir dirname

DESCRIPTION mkdir creates directory dirname.
The standard entries "." and ".." are made automatically.

FILES --

SEE ALSO rmdir to remove directories

DIAGNOSTICS "?"

BUGS No permissions are checked. The system's user ID, not that of the creator of the directory, becomes the owner of the directory.

OWNER ken, dmr

11/3/71

MKFS (I)

NAME **mkfs** -- make file system

SYNOPSIS /etc/mkfs t
 /etc/mkfs r

DESCRIPTION **mkfs** initializes either a DECtape (argument "t") or an RK03 disk pack (argument "r") so that it contains an empty file system. **mkfs** or its equivalent must be used before a tape or pack can be mounted as a file system.

In both cases the super-block, i-list, and free list are initialized, and a root directory containing entries for ". " and ".." are created. For RK03's the number of available blocks is 4872, for tapes 578.

This program is kept in /etc to avoid inadvertant use and consequent destruction of information.

FILES /dev/tap0, /dev/rk0

SEE ALSO --

DIAGNOSTICS "Arg count", "Unknown argument", "Open error".

BUGS --

OWNER ken, dmr

11/3/71

MOUNT (I)

NAME mount -- mount file system

SYNOPSIS mount special dir

DESCRIPTION mount announces to the system that a removable file system has been mounted on the device corresponding to special file special. Directory dir (which must exist already) becomes the name of the root of the newly mounted file system.

FILES --

SEE ALSO umount

DIAGNOSTICS "?" , if the special file is already in use, cannot be read, or if dir does not exist.

BUGS Should be usable only by the super-user.

OWNER ken, dmr

NAME mv -- move or rename a file

SYNOPSIS mv name₁ name₂

DESCRIPTION mv changes the name of name₁ by linking to it under the name name₂, and then unlinking name₁. Several pairs of arguments may be given. If the new name is a directory, the file is moved to that directory under its old name. Directories may only be moved within the same parent directory (just renamed).

FILES --

SEE ALSO --

DIAGNOSTICS
"?a"-- incorrect argument count
"?d"-- attempt to move a directory
"?s"-- moving file to itself
"?l"-- link error; old file doesn't exist or can't write new directory
"?u"-- can't unlink old name

BUGS If mv succeeds in removing the target file, but then is unable to link back to the old file, the result is ?l and the removal of the target file. This is common with demountable file systems and should be circumvented. Also in such cases, mv should copy if it can.

OWNER ken, dmr

NAME **nm** -- get name list

SYNOPSIS **nm** [name]

DESCRIPTION **nm** prints the symbol table from the output file of an assembler or loader run. Only relocatable, global, and undefined symbols— not absolute— are given. Each defined symbol is preceded by its value; each undefined symbol by blanks. Global symbols have their first character underlined. The output is sorted alphabetically.

If no file is given, the symbols in **a.out** are listed.

FILES **a.out**

SEE ALSO **as**, **ld**

DIAGNOSTICS "?"

BUGS --

OWNER **dmr, ken**

NAME **od** -- octal dump

SYNOPSIS **od** name [origin]

DESCRIPTION **od** dumps a file in octal, eight words per line with the origin of the line on the left. If an octal origin is given it is truncated to 0 mod 16 and dumping starts from there, otherwise from 0. Printing continues until halted by sending an interrupt signal.

FILES --

SEE ALSO db

DIAGNOSTICS "?"

BUGS Dumping does not cease at the end of the file; instead the file appears to be padded with garbage to a length of 511 mod 512 bytes.

OWNER ken, dmr

NAME `pr -- print file`

SYNOPSIS `pr [-lcm] name, ...`

DESCRIPTION `pr` produces a printed listing of one or more files. The output is separated into pages headed by the name of the file, a date, and the page number.

The optional flag -l causes each page to contain 78 lines instead of the standard 66 to accommodate legal size paper.

The optional flags -c (current date) and -m (modified date) specify which date will head all subsequent files. -m is default.

FILES /dev/ttyn to suspend messages.

SEE ALSO `cat, cp, mesg`

DIAGNOSTICS -- (files not found are ignored)

BUGS none

OWNER ken, dmr

NAME **rew** -- rewind tape

SYNOPSIS **rew** [digit]

DESCRIPTION **rew** rewinds DECtape drives. The digit is the logical tape number, and should range from 0 to 7. A missing digit indicates drive 0.

FILES /dev/tap0, ..., /dev/tap7

SEE ALSO --

DIAGNOSTICS "?" if there is no tape mounted on the indicated drive or if the file cannot be opened.

BUGS --

OWNER ken, dmr

NAME **rkd** -- dump RK disk to tape

SYNOPSIS /etc/rkd

DESCRIPTION **rkd** copies an RK03/RK05 disk pack onto nine DECtapes.

Physical I/O is done and interrupts are disabled, so time-sharing is suspended during operation of the command.

The sequence of tape drives is: 0, 1, 0, 1,

rkd exits if 0 appears in the console switches.

FILES --

SEE ALSO **rkl**

DIAGNOSTICS none; errors are retried forever

BUGS --

OWNER **ken**

NAME rkf -- format RK03 disk pack

SYNOPSIS rkf

DESCRIPTION rkf formats a virgin disk pack. Because it destroys all information on that pack, and because it is not interlocked against file system activity on the pack, the rkf program is not maintained in executable form. Instead the source must be located and assembled.

FILES none (uses physical I/O on drive 0).

SEE ALSO --

DIAGNOSTICS "error" is printed and a core image is produced if a write error occurs. A copy of the RK status register is in register 5.

BUGS As mentioned, rkf is not interlocked against system I/O; if I/O is already occurring, it will be badly disrupted. In any event, all information on the pack is destroyed.

OWNER ken, dmr

NAME `rkl -- reload RK disk from tape`

SYNOPSIS /etc/rkl

DESCRIPTION `rkl` loads an RK05/RK05 disk pack from nine DECTapes.

The program uses physical I/O with interrupts disabled; therefore time-sharing is suspended.

Only the super-user may invoke this command.

The sequence of drives is: 0, 1, 0, 1, `rkl` will cease if 0 appears in the console switches.

FILES --

SEE ALSO `rkd`

DIAGNOSTICS none; errors are retried forever

BUGS --

OWNER ken

NAME **rm** -- remove (unlink) files

SYNOPSIS **rm** name, ...

DESCRIPTION **rm** removes the entries for one or more files from a directory. If an entry was the last link to the file, the file is destroyed. Removal of a file requires write permission in its directory, but neither read nor write permission on the file itself.

Directories cannot be removed by **rm**; cf. **rmdir**.

FILES none

SEE ALSO **rmdir**, for removing directories.

DIAGNOSTICS If the file cannot be removed or does not exist, the name of the file followed by a question mark is typed.

BUGS ,
 rm probably should ask whether a read-only file is really to be removed.

OWNER ken, dmr

NAME **rmdir** -- remove directory

SYNOPSIS **rmdir** *dir*, ...

DESCRIPTION **rmdir** removes (deletes) directories. The directory must be empty (except for the standard entries ". " and "... , which **rmdir** itself removes). Write permission is required in the directory in which the directory appears.

FILES none

SEE ALSO --

DIAGNOSTICS "dir?" is printed if directory *dir* cannot be found, is not a directory, or is not removable.
"dir -- directory not empty" is printed if *dir* has entries other than ". " or "... ."

BUGS --

OWNER ken, dmr

NAME roff -- format text

SYNOPSIS roff [+number] [-number] name, ...

DESCRIPTION roff formats text according to control lines embedded in the text. The optional argument "+number" causes printing to begin at the first page with the appropriate number; "-number" causes printing to cease at the first page with a higher number.

roff is fully described in a separate publication [reference].

FILES /etc/suftab contains a list of suffixes used to guide hyphenation. /tmp/rtma, rtmb, ... temporary. /dev/ttyn to suspend messages.

SEE ALSO [reference], mesg

DIAGNOSTICS none -- files not found are ignored

BUGS roff does not check for various kinds of buffer overflow. If a fault occurs, check the input in the region where the error occurred.

OWNER jfo, dmr, ken

11/3/71

SDATE (I)

NAME **sdate** -- set date and time

SYNOPSIS **sdate** mmddhhmm

DESCRIPTION **sdate** adjusts the system's idea of the date and time. mm is the month number; dd is the day number in the month; hh is the hour number (24-hour system); mm is the minute number. For example,

sdate 10080045

sets the date to Oct. 8, 12:45 AM.

FILES none

SEE ALSO **date**

DIAGNOSTICS "?" if the date is syntactically incorrect.

BUGS none

OWNER ken, dmr

NAME `sh -- shell (command interpreter)`

SYNOPSIS `sh [name [arg1 ... [arg9]]]`

DESCRIPTION `sh` is the standard command interpreter. It is the program which reads and arranges the execution of the command lines typed by most users. It may itself be called as a command to interpret files of command lines. Before discussing the arguments to the shell used as a command, the structure of command lines themselves will be given.

Command lines are sequences of commands separated by command delimiters. Each command is a sequence of non-blank command arguments separated by blanks. The first argument specifies the name of a command to be executed. Except for certain types of special arguments discussed below, the arguments other than the command name are simply passed to the invoked command.

If the first argument represents the path name of an executable file, it is invoked; otherwise the string `"/bin/"` is prepended to the argument. (In this way the standard commands, which reside in `"/bin"`, are found.) If this search too fails a diagnostic is printed.

The remaining non-special arguments are simply passed to the command without further interpretation by the shell.

There are three command delimiters: the new line, `";"`, and `"&"`. The semicolon `";"` specifies sequential execution of the commands so separated; that is,

`coma; comb`

causes the execution first of command `coma`, then of `comb`. The ampersand `"&"` causes simultaneous execution:

`coma & comb`

causes `coma` to be called, followed immediately by `comb` without waiting for `coma` to finish. Thus `coma` and `comb` execute simultaneously. As a special case,

`coma &`

causes `coma` to be executed and the shell immediately to request another command without waiting for `coma`.

Two characters cause the immediately following string to be interpreted as a special argument to the shell itself, not passed to the command. An argument of the form "arg" causes the file arg to be used as the standard input file of the given command; an argument of the form ">arg" causes file "arg" to be used as the standard output file for the given command.

If any argument contains either of the characters "?" or "*", it is treated specially as follows. The current directory is searched for files which match the given argument. The character "*" in an argument matches any string of characters in a file name (including the null string); "?" matches any single character in a file name. Other argument characters match only the same character in the file name. For example, "*" matches all file names; "?" matches all one-character file names; "ab*.s" matches all file names beginning with "ab" and ending with ".s".

If the argument with "*" or "?" also contains a "/", a slightly different procedure is used: instead of the current directory, the directory used is the one obtained by taking the argument up to the last "/" before a "*" or "?". The matching process matches the remainder of the argument after this "/" against the files in the derived directory. For example: "/usr/dmr/a*.s" matches all files in directory "/usr/dmr" which begin with "a" and end with ".s".

In any event, a list of names is obtained which match the argument. This list is sorted into alphabetical order, and the resulting sequence of arguments replaces the single argument containing the "*" or "?". The same process is carried out for each argument with a "*" or "?" (the resulting lists are not merged) and finally the command is called with the resulting list of arguments.

For example: directory /usr/dmr contains the files a1.s, a2.s, ..., a9.s. From any directory, the command

as /usr/dmr/a?.s

calls as with arguments /usr/dmr/a1.s, /usr/dmr/a2.s, ... /usr/dmr/a9.s in that order.

The character "\ " causes the immediately following character to lose any special meaning it may have to the shell; in this way "<", ">", and other characters meaningful to the shell may be passed as part of arguments. A special case of

this feature allows the continuation of commands onto more than one line: a new-line preceded by "\ " is translated into a blank.

Sequences of characters enclosed in double ("") or single ('') quotes are also taken literally.

When the shell is invoked as a command, it has additional string processing capabilities. Recall that the form in which the shell is invoked is

```
sh [ name [ arg1 ... [ arg9 ] ] ]
```

The name is the name of a file which will be read and interpreted. If not given, this subinstance of the shell will continue to read the standard input file.

In the file, character sequences of the form "\$n", where n is a digit 0, ..., 9, are replaced by the nth argument to the invocation of the shell (arg_n). "\$0" is replaced by name.

An end-of-file in the shell's input causes it to exit. A side effect of this fact means that the way to log out from UNIX is to type an end of file.

FILES

/etc/glob

SEE ALSO

[reference], which gives the theory of operation of the shell.

DIAGNOSTICS

"?", in case of any difficulty. The most common problem is inability to find the given command. Others: input file ("<") cannot be found; no more processes can be created (this will alleviate itself with the passage of time). Note that no diagnostic is given for inability to create an output (">") file; the standard output file has already been closed when the condition is discovered and there is no place to write the diagnostic.

If a "*" or "?" is used, the glob routine is invoked; it types "No command" if it cannot find the given command, and "No match" if there were no files which matched an argument with "?" or "*".

BUGS

Better diagnostics should be provided. If a "*" or "?" is used, the command must be in /bin. (Not, for example, in the user's directory.) This is actually a glob bug.

11/3/71

SH (I)

OWNER dmr, ken

NAME **stat** -- get file status

SYNOPSIS **stat** name₁ ...

DESCRIPTION **stat** gives several kinds of information about one or more files:

i-number
access mode
number of links
owner
size in bytes
date and time of last modification
name (useful when several files are named)

All information is self-explanatory except the mode. The mode is a six-character string whose characters mean the following:

1 s: file is small (smaller than 4096 bytes)
 l: file is large

2 d: file is a directory
 x: file is executable
 u: set user ID on execution
 -: none of the above

3 r: owner can read
 -: owner cannot read

4 w: owner can write
 -: owner cannot write

5 r: non-owner can read
 -: non-owner cannot read

6 w: non-owner can write
 -: non-owner cannot write

The owner is almost always given in symbolic form; however if he cannot be found in "/etc/uids" a number is given.

If the number of arguments to **stat** is not exactly 1 a header is generated identifying the fields of the status information.

FILES /etc/uids

SEE ALSO **ls** with the "-l" option gives the same information as **stat**.

DIAGNOSTICS "name?" for any error.

BUGS none

11/3/71

STAT (I)

OWNER

dmr

NAME **strip** -- remove symbols and relocation bits

SYNOPSIS **strip** name, ...

DESCRIPTION **strip** removes the symbol table and relocation bits ordinarily attached to the output of the assembler and loader. This is useful to save space after a program has been debugged.

The effect of **strip** is the same as use of the **-s** option of **ld**.

FILES /tmp/stma, stmb ... temporary file

SEE ALSO **ld**, **as**

DIAGNOSTICS Diagnostics are given for: non-existent argument; inability to create temporary file; improper format (not an object file); inability to re-read temporary file.

BUGS --

OWNER dmr

NAME su -- become privileged user

SYNOPSIS su password

DESCRIPTION su allows one to become the super-user, who has all sorts of marvelous powers. In order for su to do its magic, the user must pass as an argument a password. If the password is correct, su will execute the shell with the UID set to that of the super-user. To restore normal UID privileges, type an end-of-file to the super-user shell.

FILES --

SEE ALSO shell

DIAGNOSTICS "Sorry" if password is wrong

BUGS --

OWNER dmr, ken

11/3/71

SUM (I)

NAME sum -- sum file

SYNOPSIS sum name

DESCRIPTION sum sums the contents of a file. In practice, it
is most often used to verify that all of a
DECtape can be read without error.

FILES none

SEE ALSO --

DIAGNOSTICS "?" if the file cannot be read at all or if an
error is discovered during the read.

BUGS none

OWNER ken

NAME `tap -- manipulate DECtape`

SYNOPSIS `tap [key] [name ...]`

DESCRIPTION `tap` saves and restores selected portions of the file system hierarchy on DECtape. Its actions are controlled by the `key` argument. The key is a string of characters containing at most one function letter and possibly one or more function modifiers. Other arguments to the command are file or directory names specifying which files are to be dumped, restored, or tabled.

The function portion of the key is specified by one of the following letters:

- r The indicated files and directories, together with all subdirectories, are dumped onto the tape. If files with the same names already exist, they are replaced (hence the "r"). "Same" is determined by string comparison, so `./abc` can never be the same as `"/usr/dmr/abc"` even if `"/usr/dmr"` is the current directory. If no file argument is given, `/` is the default.
- u updates the tape. `u` is the same as `r`, but a file is replaced only if its modification date is later than the date stored on the tape; that is to say, if it has changed since it was dumped. `u` is the default command if none is given.
- d deletes the named files and directories from the tape. At least one file argument must be given.
- x extracts the named files from the tape to the file system. The owner, mode, and date-modified are restored to what they were when the file was dumped. If no file argument is given, the entire contents of the tape are extracted.
- t lists the names of all files stored on the tape which are the same as or are hierarchically below the file arguments. If no file argument is given, the entire contents of the tape are tabled.
- l is the same as `t` except that an expanded listing is produced giving all the available information about the listed files.

The following characters may be used in addition to the letter which selects the function desired.

0, ..., 7 This modifier selects the drive on which the tape is mounted. "0" is the default.

- v Normally tap does its work silently. The v (verbose) option causes it to type the name of each file it treats preceded by a letter to indicate what is happening.

r file is being replaced
 a file is being added (not there before)
 x file is being extracted
 d file is being deleted

The v option can be used with r, u, d, and x only.

- c means a fresh dump is being created; the tape directory will be zeroed before beginning. Usable only with r and u.
- f causes new entries copied on tape to be 'fake' in that only the entries, not the data associated with the entries are updated. Such fake entries cannot be extracted. Usable only with r and u.
- w causes tap to pause before treating each file, type the indicative letter and the file name (as with v) await the user's response. Response "y" means "yes", so the file is treated. Null response means "no", and the file does not take part in whatever is being done. Response "x" means "exit"; the tap command terminates immediately. In the x function, files previously asked about have been extracted already. With r, u, and d no change has been made to the tape.
- m make (create) directories during an x if necessary.
- i ignore tape errors. It is suggested that this option be used with caution to read damaged tapes.

FILES /dev/tap0 ... /dev/tap7

SEE ALSO rk

DIAGNOSTICS RK open error
 RK read error
 RK write error
 Directory checksum
 Directory overflow

11/3/71

TAP (I)

RK overflow

Phase error (a file has changed after it was selected for dumping but before it was dumped)

BUGS

All references to "RK" should read "tape." The m option does not work correctly in all cases. The i option is not yet implemented.

OWNER

ken

NAME	tm -- provide time information		
SYNOPSIS	<u>tm</u> [command arg,]		
DESCRIPTION	<u>tm</u> is used to provide timing information. When used without an argument, output like the following is given:		
	tim	77:43:20	29.2
	ovh	13:59:42	1.2
	dsk	12:06:30	4.1
	idl	352:31:37	23.7
	usr	3:32:15	0.1
	der	5, 171	0, 0
The first column of numbers gives totals in the named categories since the last time the system was cold-booted; the second column gives the changes since the last time <u>tm</u> was invoked. The <u>tim</u> row is total real time (hours:minutes:seconds); unlike the other times, its origin is the creation date of <u>tm</u> 's temporary file. <u>ovh</u> is time spent executing in the system; <u>dsk</u> is time spent waiting for both kinds of disk I/O; <u>idl</u> is idle time; <u>usr</u> is user execution time; <u>der</u> is RF disk error count (left number) and RK disk error count (right number).			
<u>tm</u> can be invoked with arguments which are assumed to constitute a command to be timed. In this case the output is as follows:			
	tim	2.2	
	ovh	0.3	
	dsk	1.8	
	idl	0.0	
	usr	0.0	
The given times represent the number of seconds spent in each category during execution of the command.			
FILES	/tmp/ttmp, /dev/rf0 (for absolute times) contains the information used to calculate the differential times.		
SEE ALSO	format of file system (which tells where the times come from)		
DIAGNOSTICS	"?" if the command cannot be executed; "can't creat temp file" if trouble with <u>ttmp</u> ; "cant read super-block" if times cannot be read from system.		
BUGS	(1) when invoked with a command argument, everything going on at the moment is counted, not just the command itself. (2) Two users doing <u>tm</u>		

11/3/71

TM (I)

simultaneously interfere with each other's use of
the temporary file.

OWNER ken, dmr

11/3/71

TTY (I)

NAME **tty** -- get tty name

SYNOPSIS **tty**

DESCRIPTION **tty** gives the name of the user's typewriter in the form "ttyn" for n a digit. The actual path name is then "/dev/ttyn".

FILES --

SEE ALSO --

DIAGNOSTICS "not a tty" if the standard input file is not a typewriter.

BUGS --

OWNER dmr, ken

11/3/71

TYPE (I)

NAME type -- type on 2741

SYNOPSIS type name₁ ...

DESCRIPTION type produces output on an IBM 2741 terminal with a Correspondence type ball.

type uses typewriter tty5, which, because of the lack of access ports, is also used as a standard communication channel. Therefore, who should be used to verify the absence of a user on tty5.

The method is as follows: type the type command. It will wait until tty5 is dialled up. When the phone answers, depress the interrupt button after paper has been loaded, and the first file will be typed. type spaces out to the end of a sheet of paper and waits until the interrupt button is depressed before beginning each new file.

FILES /dev/tty5

SEE ALSO who

DIAGNOSTICS --

BUGS Obviously some scheme is needed to prevent interference between normal users and type. The best thing would be to support 2741's as a standard terminal.

OWNER dmr

11/3/71

UMOUNT (I)

NAME **umount** -- dismount file system

SYNOPSIS **umount** special

DESCRIPTION **umount** announces to the system that the removable file system previously mounted on special file special is to be removed.

Only the super-user may issue this command.

FILES --

SEE ALSO **mount**

DIAGNOSTICS ?

BUGS This command should be restricted to the super-user.

OWNER ken, dmr

NAME **un** -- undefined symbols

SYNOPSIS **un** [name]

DESCRIPTION **un** prints a list of undefined symbols from an assembly or loader run. If the file argument is not specified, a.out is the default. Names are listed alphabetically except that non-global symbols come first. Undefined global symbols (unresolved external references) have their first character underlined.

FILES **a.out**

SEE ALSO **as, ld**

DIAGNOSTICS "?" if the file cannot be found.

BUGS --

OWNER **dmr, ken**

NAME **wc** -- get (English) word count

SYNOPSIS **wc name₁ ...**

DESCRIPTION **wc** provides a count of the words, text lines, and roff control lines for each argument file.

A text line is a sequence of characters not beginning with **"."** and ended by a new-line. A roff control line is a line beginning with **". ."**. A word is a sequence of characters bounded by the beginning of a line, by the end of a line, or by a blank or a tab.

FILES --

SEE ALSO roff

DIAGNOSTICS none; arguments not found are ignored.

BUGS --

OWNER jfo

11/3/71

WHO (I)

NAME **who** -- who is on the system

SYNOPSIS **who**

DESCRIPTION **who** lists the name, typewriter channel, and login time for each current UNIX user.

FILES /tmp/utmp contains the necessary information; it is maintained by init.

SEE ALSO /etc/init

DIAGNOSTICS --

BUGS --

OWNER dmr, ken

NAME write -- write to another user

SYNOPSIS write user

DESCRIPTION write copies lines from your typewriter to that of another user. When first called, write sends the message
 message from yourname...

The recipient of the message should write back at this point. Communication continues until an end of file is read from the typewriter or an interrupt is sent. At that point write writes "EOT" on the other terminal.

Permission to write may be denied or granted by use of the mesq command. At the outset writing is allowed. Certain commands, in particular roff and pr, disallow messages in order to prevent messy output.

If the character "!" is found at the beginning of a line, write calls the mini-shell msh to execute the rest of the line as a command.

The following protocol is suggested for using write: When you first write to another user, wait for him to write back before starting to send. Each party should end each message with a distinctive signal ("(o)" for "over" is conventional) that the other may reply. "(oo)" (for "over and out") is suggested when conversation is about to be terminated.

FILES /tmp/utmp is used to discover the target user's typewriter channel and the sending user's name. msh is used to execute commands.

SEE ALSO mesq

DIAGNOSTICS "user not logged in"; "permission denied".

BUGS --

OWNER dmr, ken

11/3/71

SYS BREAK (II)

NAME break -- set program break

SYNOPSIS sys break; addr / break = 17.

DESCRIPTION break sets the system's idea of the highest location used by the program to addr. Locations greater than addr and below the stack pointer are not swapped and are thus liable to unexpected modification.

If the argument is 0 or higher than the stack pointer the entire 4K word user core area is swapped.

When a program begins execution via exec the break is set at the highest location defined by the program and data storage areas. Ordinarily, therefore, only programs with growing data areas need to use break.

FILES --

SEE ALSO exec

DIAGNOSTICS none; strange addresses cause the break to be set to include all of core.

BUGS --

OWNER ken, dmr

11/3/71

SYS CEMT (II)

NAME cemt -- catch emt traps

SYNOPSIS sys cemt; arg / cemt = 29.; not in assembler

DESCRIPTION This call allows one to catch traps resulting from the emt instruction. Arg is a location within the program; emt traps are sent to that location. The normal effect of emt traps may be restored by giving an arg equal to 0.

Prior to the use of this call, the result of an emt instruction is a simulated rts instruction. The operand field is interpreted as a register, and an rts instruction is simulated for that register (after verifying that various registers have appropriate values). This feature is useful for debugging, since the most dangerous program bugs usually involve an rts with bad data on the stack or in a register.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME chdir -- change working directory

SYNOPSIS sys chdir; dirname / chdir = 12.

DESCRIPTION dirname is address of the pathname of a directory, terminated by a 0 byte. chdir causes this directory to become the current working directory.

FILES --

SEE ALSO --

DIAGNOSTICS The error bit (c-bit) is set if the given name is not that of a directory.

BUGS --

OWNER ken, dmr

NAME chmod -- change mode of file

SYNOPSIS sys chmod; name; mode / chmod = 15.

DESCRIPTION The file whose name is given as the null-terminated string pointed to by name has its mode changed to mode. Modes are constructed by oring together some combination of the following:

01 write, non-owner
02 read, non-owner
04 write, owner
10 read, owner
20 executable
40 set user ID on execution

Only the owner of a file (or the super-user) may change the mode.

FILES --

SEE ALSO --

DIAGNOSTICS Error bit (c-bit) set if name cannot be found or if current user is neither the owner of the file nor the super-user.

BUGS --

OWNER ken, dmr

11/3/71

SYS CHOWN (II)

NAME chown -- change owner of file

SYNOPSIS sys chown; name; owner / chown = 16.

DESCRIPTION The file whose name is given by the null-terminated string pointed to by name has its owner changed to owner. Only the present owner of a file (or the super-user) may donate the file to another user. Also, one may not change the owner of a file with the set-user-ID bit on, otherwise one could create Trojan Horses able to misuse other's files.

FILES --

SEE ALSO /etc/uids has the mapping between user names and user numbers.

DIAGNOSTICS The error bit (c-bit) is set on illegal owner changes.

BUGS --

OWNER ken, dmr

11/3/71

SYS CLOSE (II)

NAME close '-- close a file

SYNOPSIS (file descriptor in r0)
 sys close / close = 6.

DESCRIPTION Given a file descriptor such as returned from an
 open or creat call, close closes the associated
 file. A close of all files is automatic on exit,
 but since processes are limited to 10 simultane-
 ously open files, close is necessary to programs
 which deal with many files.

FILES --

SEE ALSO creat, open

DIAGNOSTICS The error bit (c-bit) is set for an unknown file
 descriptor.

BUGS --

OWNER ken, dmr

NAME `creat` -- create a new file

SYNOPSIS `sys creat; name; mode` / `creat = 8.`
 (file descriptor in `r0`)

DESCRIPTION creat creates a new file or prepares to rewrite
 an existing file called name; name is the address
 of a null-terminated string. If the file did not
 exist, it is given mode mode; if it did exist,
 its mode and owner remain unchanged but it is
 truncated to 0 length.

The file is also opened for writing, and its file
descriptor is returned in `r0`.

The mode given is arbitrary; it need not allow
writing. This feature is used by programs which
deal with temporary files of fixed names. The
creation is done with a mode that forbids writ-
ing. Then if a second instance of the program
attempts a creat, an error is returned and the
program knows that the name is unusable for the
moment.

If the last link to an open file is removed, the
file is not destroyed until the file is closed.

FILES --

SEE ALSO `write, close`

DIAGNOSTICS The error bit (c-bit) may be set if: a needed
 directory is not readable; the file does not
 exist and the directory in which it is to be
 created is not writable; the file does exist and
 is unwritable; the file is a directory.

BUGS --

OWNER ken, dmr

NAME	<code>exec -- execute a file</code>	
SYNOPSIS	<code>sys exec; name; args</code> / exec = 11. ... <code>name: <...\\0></code> ... <code>args: arg1; arg2; ...; 0</code> <code>arg1: <...\\0></code> ...	
DESCRIPTION	<p><code>exec</code> overlays the calling process with the named file, then transfers to the beginning of the core image of the file. The first argument to <code>exec</code> is a pointer to the name of the file to be executed. The second is the address of a list of pointers to arguments to be passed to the file. Conventionally, the first argument is the name of the file. Each pointer addresses a string terminated by a null byte.</p> <p>There can be no return from the file; the calling core image is lost.</p> <p>The program break is set from the executed file; see the format of <code>a.out</code>.</p> <p>Once the called file starts execution, the arguments are passed as follows. The stack pointer points to the number of arguments. Just above this number is a list of pointers to the argument strings.</p> <pre> sp-> nargs arg1 ... argn arg1: <arg1\\0> ... argn: <argn\\0> </pre> <p>The arguments are placed as high as possible in core: just below 60000(8).</p> <p>Files remain open across <code>exec</code> calls. However, the illegal instruction, <code>emt</code>, <code>quit</code>, and interrupt trap specifications are reset to the standard values. (See <u>ilqins</u>, <u>cemt</u>, <u>quit</u>, <u>intr</u>.)</p> <p>Each user has a <u>real</u> user ID and an <u>effective</u> (The real ID identifies the person using the system; the effective ID determines his access privileges.) <code>exec</code> changes the effective user ID to the owner of the executed file if the file has the "set-user-ID" mode. The real user ID is not affected.</p>	

11/3/71

SYS EXEC (II)

FILES --

SEE ALSO fork

DIAGNOSTICS If the file cannot be read or if it is not executable, a return from exec constitutes the diagnostic. The error bit (c-bit) is set.

BUGS --

OWNER ken, dmr

11/3/71

SYS EXIT (II)

NAME **exit** -- terminate process

SYNOPSIS **sys exit / exit = 1**

DESCRIPTION **exit** is the normal means of terminating a process. All files are closed and the parent process is notified if it is executing a **wait**.
This call can never return.

FILES --

SEE ALSO **sys wait**

DIAGNOSTICS -

BUGS --

OWNER ken, dmr

NAME **fork** -- spawn new process

SYNOPSIS **sys fork / fork = 2.**
 (**new process return**)
 (**old process return**)

DESCRIPTION **fork** is the only way new processes are created.
The new process's core image is a copy of that of
the caller of **fork**; the only distinction is the
return location and the fact that r0 in the old
process contains the process ID of the new pro-
cess. This process ID is used by wait.

FILES --

SEE ALSO **sys wait, sys exec**

DIAGNOSTICS The error bit (c-bit) is set in the old process
if a new process could not be created because of
lack of swap space.

BUGS See wait for a subtle bug in process destruction.

OWNER ken, dmr

11/3/71

SYS FSTAT (II)

NAME fstat -- get status of open file

SYNOPSIS (file descriptor in r0)
 sys fstat; buf / fstat = 28.

DESCRIPTION This call is identical to stat, except that it
 operates on open files instead of files given by
 name. It is most often used to get the status of
 the standard input and output files, whose names
 are unknown.

FILES --

SEE ALSO sys stat

DIAGNOSTICS The error bit (c-bit) is set if the file descrip-
 tor is unknown.

BUGS --

OWNER ken, dmr

NAME **getuid** -- get user identification

SYNOPSIS **sys getuid / getuid = 24.**
 (user ID in r0)

DESCRIPTION **getuid** returns the real user ID of the current process. The real user ID identifies the person who is logged in, in contradistinction to the effective user ID, which determines his access permission at each moment. It is thus useful to programs which operate using the "set user ID" mode, to find out who invoked them.

FILES /etc/uids can be used to map the user ID number into a name.

SEE ALSO **setuid**

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

11/3/71

SYS GTTY (II)

NAME **gtty** -- get typewriter status

SYNOPSIS (file descriptor in r0)
 sys **gtty**; arg / **gtty** = 32.; not in assembler
 ...
 arg: .=.+6

DESCRIPTION **gtty** stores in the three words addressed by arg
 the status of the typewriter whose file descriptor
 is given in r0. The format is the same as
 that passed by stty.

FILES --

SEE ALSO **stty**

DIAGNOSTICS Error bit (c-bit) is set if the file descriptor
 does not refer to a typewriter.

BUGS --

OWNER ken, dmr

NAME ilgins -- catch illegal instruction trap

SYNOPSIS sys ilgins; arg / ilgins = 33.; not in assembler

DESCRIPTION ilgins allows a program to catch illegal instruction traps. If arg is zero, the normal instruction trap handling is done: the process is terminated and a core image is produced. If arg is a location within the program, control is passed to arg when the trap occurs.

This call is used to implement the floating point simulator, which catches and interprets 11/45 floating point instructions.

FILES --

SEE ALSO fptrap, the floating point package

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME intr -- set interrupt handling

SYNOPSIS sys intr; arg / intr = 27.

DESCRIPTION When arg is 0, interrupts (ASCII DELETE) are ignored. When arg is 1, interrupts cause their normal result, that is, force an exit. When arg is a location within the program, control is transferred to that location when an interrupt occurs.

After an interrupt is caught, it is possible to resume execution by means of an rti instruction; however, great care must be exercised, since all I/O is terminated abruptly upon an interrupt. In particular, reads of the typewriter tend to return with 0 characters read, thus simulating an end of file.

FILES --

SEE ALSO quit

DIAGNOSTICS --

BUGS It should be easier to resume after an interrupt, but I don't know how to make it work.

OWNER ken, dmr

11/3/71

SYS LINK (1)

NAME link -- link to a file

SYNOPSIS sys link; name₁; name₂ / link = 9.

DESCRIPTION A link to name₁ is created; the link has name name₂. Either name may be an arbitrary path name.

FILES --

SEE ALSO unlink

DIAGNOSTICS The error bit (c-bit) is set when name₁ cannot be found; when name₂ already exists; when the directory of name₂ cannot be written; when an attempt is made to link to a directory by a user other than the super-user.

BUGS --

OWNER ken, dmr

NAME mkdir -- make a directory

SYNOPSIS sys mkdir; name; mode / mkdir = 14.

DESCRIPTION mkdir creates an empty directory whose name is the null-terminated string pointed to by name. The mode of the directory is mode. The special entries ". " and ". ." are not present.
mkdir can only be invoked by the super-user.

FILES --

SEE ALSO mkdir command

DIAGNOSTICS Error bit (c-bit) is set if the directory already exists or if the user is not the super-user.

BUGS --

OWNER ken, dmr

NAME mount -- mount file system

SYNOPSIS sys mount; special; name / mount = 21.; not in assembler

DESCRIPTION mount announces to the system that a removable file system has been mounted on special file special; from now on, references to file name will refer to the root file on the newly mounted file system. Special and name are pointers to null-terminated strings containing the appropriate path names.

Name must exist already. If it had useful contents, they are inaccessible while the file system is mounted.

Almost always, name should be a directory so that an entire file system, not just one file, may exist on the removable device.

FILES --

SEE ALSO umount

DIAGNOSTICS Error bit (c-bit) set if special is inaccessible or dir does not exist.

BUGS At most one removable device can be mounted at a time. The use of this call should be restricted to the super-user.

OWNER ken, dmr

NAME `open -- open for reading or writing`

SYNOPSIS `sys open; name; mode / open = 5.`
 `(descriptor in r0)`

DESCRIPTION `open` opens the file `name` for reading (if `mode` is 0) or writing (if `mode` is non-zero). `name` is the address of a string of ASCII characters representing a path name, terminated by a null character.

The file descriptor should be saved for subsequent calls to read (or write) and close.

In both the read and write case the file pointer is set to the beginning of the file.

If the last link to an open file is removed, the file is not destroyed until it is closed.

FILES --

SEE ALSO `creat, read, write, close`

DIAGNOSTICS The error bit (c-bit) is set if the file does not exist, if one of the necessary directories does not exist or is unreadable, or if the file is not readable.

BUGS --

OWNER ken, dmr

11/3/71

SYS QUIT (II)

NAME quit -- turn off quit signal

SYNOPSIS sys quit; flag / quit = 26.

DESCRIPTION When flag is 0, this call disables quit signals from the typewriter (ASCII FS). When flag is 1, quits are re-enabled, and cause execution to cease and a core image to be produced. When flag is an address in the program, a quit causes control to be sent to that address.

Quits should be turned off only with due consideration.

FILES --

SEE ALSO sys intr turns off interrupts

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME read -- read from file

SYNOPSIS (file descriptor in r0)
 sys read; buffer; nchars / read = 3.
 (nread in r0)

DESCRIPTION A file descriptor is a word returned from a successful open call.

Buffer is the location of nchars contiguous bytes into which the input will be placed. It is not guaranteed that all nchars bytes will be read, however; for example if the file refers to a typewriter at most one line will be returned. In any event the number of characters read is returned in r0.

 If r0 returns with value 0, then end-of-file has been reached.

FILES --

SEE ALSO open

DIAGNOSTICS As mentioned, r0 is 0 on return when the end of the file has been reached. If the read was otherwise unsuccessful the error bit (c-bit) is set. Many conditions, all rare, can generate an error: physical I/O errors, bad buffer address, preposterous nchars, file descriptor not that of an input file.

BUGS --

OWNER ken, dmr

NAME rele -- release processor

SYNOPSIS sys rele / rele = 0; not in assembler

DESCRIPTION This call causes the process to be swapped out immediately if another process wants to run. Its main reason for being is internal to the system, namely to implement timer-runout swaps. However, it can be used beneficially by programs which wish to loop for some reason without consuming more processor time than necessary.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME seek -- move read/write pointer

SYNOPSIS (file descriptor in r0)
 sys seek; offset; ptrname / seek = 19.

DESCRIPTION The file descriptor refers to a file open for reading or writing. The read (or write) pointer for the file is set as follows:

if ptrname is 0, the pointer is set to offset.

if ptrname is 1, the pointer is set to its current location plus offset.

if ptrname is 2, the pointer is set to the size of the file plus offset.

FILES --

SEE ALSO tell

DIAGNOSTICS The error bit (c-bit) is set for an undefined file descriptor.

BUGS A file can conceptually be as large as $2^{**}20$ bytes. Clearly only $2^{**}16$ bytes can be addressed by seek. The problem is most acute on the tape files and RK and RF. Something is going to be done about this.

OWNER ken, dmr

11/3/71

SYS SETUID (II)

NAME **setuid** -- set process ID

SYNOPSIS (process ID in r0)
 sys setuid / setuid = 23.

DESCRIPTION The user ID of the current process is set to the argument in r0. Both the effective and the real user ID are set. This call is only permitted to the super-user.

FILES --

SEE ALSO **getuid**

DIAGNOSTICS Error bit (c-bit) is set if the current user ID is not that of the super-user.

BUGS --

OWNER ken, dmr

NAME smdate -- set modified date on file

SYNOPSIS (time to AC-MQ)
 sys smdate; file / smdate = 30.; not in assembler

DESCRIPTION File is the address of a null-terminated string giving the name of a file. The modified time of the file is set to the time given in the AC-MQ registers.

This call is allowed only to the super-user.

FILES --

SEE ALSO --

DIAGNOSTICS Error bit is set if the user is not the super-user or if the file cannot be found.

BUGS --

OWNER ken, dmr

NAME stat -- get file status
SYNOPSIS sys stat; name; buf / stat = 18.
DESCRIPTION name points to a null-terminated string naming a file; buf is the address of a 34(10) byte buffer into which information is placed concerning the file. It is unnecessary to have any permissions at all with respect to the file, but all directories leading to the file must be readable.

After stat, buf has the following format:

buf, +1	i-number
+2,+3	flags (see below)
+4	number of links
+5	user ID of owner
+6,+7	size in bytes
+8,+9	first indirect block or contents block
...	
+22,+23	eighth indirect block or contents block
+24,+25,+26,+27	creation time
+28,+29,+30,+31	modification time
+32,+33	unused

The flags are as follows:

100000	used (always on)
040000	directory
020000	file has been modified (always on)
010000	large file
000040	set user ID
000020	executable
000010	read, owner
000004	write, owner
000002	read, non-owner
000001	write, non-owner

FILES --
SEE ALSO fstat
DIAGNOSTICS Error bit (c-bit) is set if the file cannot be found.
BUGS The format is going to change someday.
OWNER ken, dmr

11/3/71

SYS STIME (II)

NAME stime -- set time

SYNOPSIS (time in AC-MQ)
sys stime / stime = 25.; not in assembler

DESCRIPTION stime sets the system's idea of the time and date. Only the super-user may use this call.

FILES --

SEE ALSO sys time

DIAGNOSTICS Error bit (c-bit) set if user is not the super-user.

BUGS --

OWNER ken, dmr

NAME **stty** -- set mode of typewriter

SYNOPSIS (file descriptor in r0)
 sys stty; arg / stty = 31.; not in assembler
 ...
arg: dcrsr; dcpsr; mode

DESCRIPTION **stty** sets mode bits for a typewriter whose file descriptor is passed in r0. First, the system delays until the typewriter is quiescent. Then, the argument dcrsr is placed into the typewriter's reader control and status register, and dcpsr is placed in the printer control and status register. The DC-11 manual must be consulted for the format of these words. For the purpose of this call, the most important rôle of these arguments is to adjust to the speed of the typewriter.

The mode arguments contains several bits which determine the system's treatment of the typewriter:

200	even (M37 tty) parity allowed
100	odd (non-M37 tty) allowed
040	raw mode: wake up on all characters
020	map CR into LF; echo LF or CR as CR-LF
010	don't echo (half duplex)
004	map upper case to lower case on input (M33 TTY)

Characters with the wrong parity, as determined by bits 200 and 100, are ignored.

In raw mode, every character is passed back immediately to the program. No erase or kill processing is done; the end-of-file character (EOT), the interrupt character (DELETE) and the quit character (FS) are not treated specially.

Mode 020 causes input carriage returns to be turned into new-lines; input of either CR or LF causes CR-LF both to be echoed (used for GE Terminate 300's).

FILES --

SEE ALSO **gtty**

DIAGNOSTICS The error bit (c-bit) is set if the file descriptor does not refer to a typewriter.

BUGS This call should be used with care. It is all too easy to turn off your typewriter.

OWNER ken, dmr

NAME tell -- get file pointer

SYNOPSIS (file descriptor in r0)
sys tell; offset; ptrname / tell = 20.
(value returned in r0)

DESCRIPTION The file descriptor refers to an open file. The value returned in r0 is one of:

if ptrname is 0, the value returned is offset;

if ptrname is 1, the value is the current pointer plus offset;

if ptrname is 2, the value returned is the number of bytes in the file plus offset.

FILES --

SEE ALSO seek

DIAGNOSTICS The error bit (c-bit) is set if the file descriptor is unknown.

BUGS Tell doesn't work. Complain if you need it.

OWNER ken, dmr

11/3/71

SYS TIME (19)

NAME time -- get time of year

SYNOPSIS sys time / time = 13.
 (time AC-MQ)

DESCRIPTION time returns the time since 00:00:00, Jan. 1,
 1971, measured in sixtieths of a second. The
 high order word is in the AC register and the low
 order is in the MQ.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS The chronological-minded user will note that
 2^{32} sixtieths of a second is only about 2.5
 years.

OWNER ken, dmr

NAME umount -- dismount file system

SYNOPSIS sys umount; special / umount = 22.; not in assembler

DESCRIPTION umount announces to the system that special file special is no longer to contain a removable file system. The file associated with the special file reverts to its ordinary interpretation (see mount).

The user must take care that all activity on the file system has ceased.

FILES --

SEE ALSO mount

DIAGNOSTICS Error bit (c-bit) set if no file system was mounted on the special file.

BUGS Use of this call should be restricted to the super-user.

OWNER ken, dmr

11/3/71

SYS UNLINK (II)

NAME unlink -- remove directory entry

SYNOPSIS sys unlink; name / unlink = 10.

DESCRIPTION Name points to a null-terminated string. Unlink removes the entry for the file pointed to by name from its directory. If this entry was the last link to the file, the contents of the file are freed and the file is destroyed. If, however, the file was open in any process, the actual destruction is delayed until it is closed, even though the directory entry has disappeared.

FILES --

SEE ALSO link

DIAGNOSTICS The error bit (c-bit) is set to indicate that the file does not exist or that its directory cannot be written. Write permission is not required on the file itself. It is also illegal to unlink a directory (except for the super-user).

BUGS Probably write permission should be required to remove the last link to a file, but this gets in other problems (namely, one can donate an undeletable file to someone else).

If the system crashes while a file is waiting to be deleted because it is open, the space is lost.

OWNER ken, dmr

11/3/71

SYS WAIT (II)

NAME	wait -- wait for process to die
SYNOPSIS	sys wait / wait = 7. (process ID in r0)
DESCRIPTION	<u>wait</u> causes its caller to delay until one of its child processes terminates. If any child has already died, return is immediate; if there are no children, return is immediate with the error bit set. In the case of several children several <u>waits</u> are needed to learn of all the deaths.
FILES	--
SEE ALSO	fork
DIAGNOSTICS	error bit (c-bit) on if no children not previously waited for.
BUGS	A child which dies but is never waited for is not really gone in that it still consumes disk swap and system table space. This can make it impossible to create new processes. The bug can be noticed when several "&" separators are given to the shell not followed by a command without an ampersand. Ordinarily things clean themselves up when an ordinary command is typed, but it is possible to get into a situation in which no commands are accepted, so no <u>waits</u> are done; the system is then hung. The fix, probably, is to have a new kind of <u>fork</u> which creates a process for which no <u>wait</u> is necessary (or possible); also to limit the number of active or inactive descendants allowed to a process.
OWNER	ken. dmr

NAME write -- write on file

SYNOPSIS (file descriptor in r0)
 sys write; buffer; nchars / write = 4.
 (number written in r0)

DESCRIPTION A file descriptor is a word returned from a successful open or creat call.

buffer is the address of nchars contiguous bytes which are written on the output file. The number of characters actually written is returned in r0. It should be regarded as an error if this is not the same as requested.

For disk and tape files, writes which are multiples of 512 characters long and begin on a 512-byte boundary are more efficient than any others.

FILES --

SEE ALSO sys creat, sys open

DIAGNOSTICS The error bit (c-bit) is set on an error: bad descriptor, buffer address, or count. physical I/O errors;

BUGS --

OWNER ken, dmr

NAME atof -- ascii to floating

SYNOPSIS jsr r5,atof; subr

DESCRIPTION atof will convert an ascii stream to a floating number returned in fr0. The subroutine subr is called on r5 for each character of the ascii stream. subr should return the character in r0. The first character not used in the conversion is left in r0. The floating point simulation should be active in either floating or double mode, but in single precision integer mode.

FILES kept in /etc/liba.a

SEE ALSO fptrap

DIAGNOSTICS --

BUGS The subroutine subr should not disturb any registers.

OWNER ken

11/3/71

ATOI (III)

NAME atoi -- ascii to integer

SYNOPSIS jsr r5,atoi; subr

DESCRIPTION atoi will convert an ascii stream to a binary number returned in mq. The subroutine subr is called on r5 for each character of the ascii stream. subr should return the character in r0. The first character not used in the conversion is left in r0.

FILES kept in /etc/liba.a

SEE ALSO --

DIAGNOSTICS --

BUGS The subroutine subr should not disturb any registers.

OWNER ken

11/3/71

CTIME (III)

NAME ctime -- convert date and time to ASCII

SYNOPSIS (move time to AC-MQ)
mov \$buffer,r0
jsr pc,ctime

DESCRIPTION The buffer is 15 characters long. The time has the format

Oct 9 17:32:24

The input time is in the AC and MQ registers in the form returned by sys time.

FILES kept in /etc/liba.a

SEE ALSO ptime, to print time; sys time

DIAGNOSTICS --

BUGS The time is not taken modulo 1 year. (Jan 1 comes out Dec 32.) Also, the clock period is only a couple of years.

OWNER dmr

11/3/71

EXP (III)

NAME exp -- exponential function

SYNOPSIS jsr r5,exp

DESCRIPTION The exponential of fr0 is returned in fr0. The floating point simulation should be active in either floating or double mode, but in single precision integer mode.

FILES kept in /etc/liba.a

SEE ALSO fptrap

DIAGNOSTICS --

BUGS Large arguments will cause an overflow fault from the floating point simulator.

OWNER ken

NAME fptrap -- floating point simulator

SYNOPSIS sys 33.; fptrap

DESCRIPTION fptrap is a program designed to pick up illegal instruction in order to simulate a sub-set of the 11/45 floating point hardware.

FILES kept in /etc/liba.a

SEE ALSO as, PDP-11/45 manual

DIAGNOSTICS none, hardware gives no diagnostics.

BUGS The simulation, if unsuccessful for any reason gives an IOT fault from inside the simulator. This should be handled better.

OWNER ken, dmr

NAME ftoa -- floating to ascii conversion

SYNOPSIS jsr r5,ftoa; subr

DESCRIPTION ftoa will convert the floating point number in fr0 into ascii in the form [-]d.dddddddde[-]dd*. The floating point simulator should be active in either floating or double mode, but in single integer mode. For each character generated by ftoa, the subroutine subr is called on register r5 with the character in r0.

FILES kept in /etc/liba.a

SEE ALSO fptrap

DIAGNOSTICS --

BUGS The subroutine subr should not disturb any registers.

OWNER ken

NAME `getw, getc, fopen -- buffered input`

SYNOPSIS

```

    mov      $filename,r0
    jsr      r5,fopen; iobuf

    jsr      r5,getc; iobuf
    (character in r0)

    jsr      r5,getw; iobuf
    (word in r0)

```

DESCRIPTION

These routines are used to provide a buffered input facility. iobuf is the address of a 134(10) byte buffer area whose contents are maintained by these routines. Its format is:

```

ioptr: .=.+2           / file descriptor
      .=.+2           / characters left in buffer
      .=.+2           / ptr to next character
      .=.+128.        / the buffer

```

fopen should be called initially to open the file. On return, the error bit (c-bit) is set if the open failed. If fopen is never called, get will read from the standard input file.

getc returns the next byte from the file in r0. The error bit is set on end of file or a read error.

getw returns the next word in r0. getc and getw may be used alternately; there are no odd/even problems.

iobuf must be provided by the user; it must be on a word boundary.

FILES kept in /etc/liba.a

SEE ALSO

sys open, sys read; putc, putw, fcreat

DIAGNOSTICS

c-bit set on EOF or error

BUGS

for greater speed, the buffer should be 512 bytes long. Unfortunately, this will cause several existing programs to stop working.

OWNER

dmr

NAME itoa -- integer to ascii conversion

SYNOPSIS jsr r5,itoa; subr

DESCRIPTION itoa will convert the number in r0 into ascii decimal possibly preceded by a - sign. For each character generated by itoa, the subroutine subr is called on register r5 with the character in r0.

FILES kept in /etc/liba.a

SEE ALSO --

DIAGNOSTICS --

BUGS The subroutine subr should not disturb any registers.

OWNER ken

NAME log -- logarithm base e
SYNOPSIS jsr r5,log
DESCRIPTION The logarithm base e of fr0 is returned in fr0.
The floating point simulation should be active in
either floating or double mode, but in single
precision integer mode.
FILES kept in /etc/liba.a
SEE ALSO fptrap
DIAGNOSTICS The error bit (c-bit) is set if the input argu-
ment is less than or equal to zero.
BUGS --
OWNER ken

NAME `mesg -- write message on typewriter`

SYNOPSIS `jsr r5,mesg; <Now is the time\0>; .even`

DESCRIPTION mesq writes the string immediately following its call onto the standard output file. The string is terminated by a 0 byte.

FILES kept in /etc/liba.a, standard output file

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME ptime -- print date and time

SYNOPSIS (move time to ac-mq)
mov file,r0
jsr pc,ptime

DESCRIPTION ptime prints the date and time in the form
Oct 9 17:20:33
on the file whose file descriptor is in r0. The string is 15 characters long. The time to be printed is placed in the AC and MQ registers in the form returned by sys time.

FILES kept in /etc/liba.a

SEE ALSO sys time, ctime (used to do the conversion)

DIAGNOSTICS --

BUGS see ctime

OWNER dmr, ken

11/3/71

PUTC, PUTW, FCREAT, FLUSH (III)

NAME	putc, putw, fcreat, flush -- buffered output
SYNOPSIS	<pre>mov \$filename,r0 jsr r5,fcreat; iobuf (get byte in r0) jsr r5,putc; iobuf (get word in r0) jsr r5,putw; iobuf jsr r5,flush; iobuf</pre>
DESCRIPTION	<p><u>fcreat</u> creates the given file (mode 17) and sets up the buffer <u>iobuf</u> (size 134(10) bytes); <u>putc</u> and <u>putw</u> write a byte or word respectively onto the file; <u>flush</u> forces the contents of the buffer to be written, but does not close the file. The format of the buffer is:</p> <pre>iobuf: .=.+2 / file descriptor .=.+2 / characters unused in buffer .=.+2 / ptr to next free character .=.+128. / buffer</pre> <p><u>fcreat</u> sets the error bit (c-bit) if the file creation failed; none of the other routines return error information.</p> <p>Before terminating, a program should call <u>flush</u> to force out the last of the output.</p> <p>The user must supply <u>iobuf</u>, which should begin on a word boundary.</p>
FILES	kept in /etc/liba.a
SEE ALSO	sys creat; sys write; getc, getw, fopen
DIAGNOSTICS	error bit possible on <u>fcreat</u> call
BUGS	buffers should be changed to 512 bytes.
OWNER	dmr

NAME sin, cos -- sine cosine
SYNOPSIS jsr r5,sin (cos)
DESCRIPTION The sine (cosine) of fr0 (radians) is returned in
 fr0. The floating point simulation should be
 active in either floating or double mode, but in
 single precision integer mode. All floating
 registers are used.
FILES kept in /etc/liba.a
SEE ALSO fptrap
DIAGNOSTICS --
BUGS Size of the argument should be checked to make
 sure the result is meaningful.
OWNER ken, dmr

NAME **switch** -- switch on value

SYNOPSIS (switch value in r0)
 jsr r5,switch; swtab
 (not-found return)
 ...
 swtab: val1; lab1;
 ...
 valn; labn
 ..; 0

DESCRIPTION **switch** compares the value of r0 against each of the val_i; if a match is found, control is transferred to the corresponding lab_i (after popping the stack once). If no match has been found by the time a null lab_i occurs, **switch** returns.

FILES kept in /etc/liba.a

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

11/3/71

/DEV/MEM (IV)

NAME mem -- core memory

SYNOPSIS --

DESCRIPTION mem maps the core memory of the computer into a file. It may be used, for example, to examine, and even to patch the system using the debugger.

Mem is a byte-oriented file; its bytes are numbered 0 to 65,535.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS If a location not corresponding to implemented memory is read or written, the system will incur a bus-error trap and, in panic, will reboot itself.

OWNER ken, dmr

NAME ppt -- punched paper tape

SYNOPSIS --

DESCRIPTION ppt refers to the paper tape reader or punch, depending on whether it is read or written.
When ppt is opened for writing, a 100-character leader is punched. Thereafter each byte written is punched on the tape. No editing of the characters is performed. When the file is closed, a 100-character trailer is punched.
When ppt is opened for reading, the process waits until tape is placed in the reader and the reader is on-line. Then requests to read cause the characters read to be passed back to the program, again without any editing. This means that several null characters will usually appear at the beginning of the file; they correspond to the tape leader. Likewise several nulls are likely to appear at the end. End-of-file is generated when the tape runs out.
Seek calls for this file are meaningless and are effectively ignored (however, the read/write pointers are maintained and an arbitrary sequence of reads or writes intermixed with seeks will give apparently correct results when checked with tell).

FILES --

SEE ALSO lbppt, dbppt, bppt format

DIAGNOSTICS --

BUGS Previously, there were separate special files for ASCII tape (which caused null characters to be suppressed) and binary tape (which used a blocked format with checksums). These notions were conceptually quite attractive, but they were discarded to save space in the system.

OWNER ken, dmr

NAME rf0 -- RF11-RS11 fixed-head disk file

SYNOPSIS --

DESCRIPTION This file refers to the entire RF disk. It may be either read or written, although writing is inherently very dangerous, since a file system resides there.

The disk contains 1024 256-word blocks, numbered 0 to 1023. Like the other block-structured devices (tape, RK disk) this file is addressed in blocks, not bytes. This has two consequences: seek calls refer to block numbers, not byte numbers; and sequential reading or writing always advance the read or write pointer by at least one block. Thus successive reads of 10 characters from this file actually read the first 10 characters from successive blocks.

FILES --

SEE ALSO /dev/tap0, /dev/rk0

DIAGNOSTICS --

BUGS The fact that this device is addressed in terms of blocks, not bytes, is extremely unfortunate. It is due entirely to the fact that read and write pointers (and consequently the arguments to seek and tell) are single-precision numbers. This really has to be changed but unfortunately the repercussions are serious.

OWNER ken, dmr

11/3/71

/DEV/RKO (IV)

NAME rk0 -- RK03 (or RK05) disk .

SYNOPSIS --

DESCRIPTION rk0 refers to the entire RK03 disk as a single sequentially-addressed file. Its 256-word blocks are numbered 0 to 4871. Like the RF disk and the tape files, its addressing is block-oriented. Consult the /dev/rf0 section.

FILES --

SEE ALSO /dev/rf0, /dev/tap0

DIAGNOSTICS --

BUGS See /dev/rf0

OWNER ken, dmr

11/3/71

/DEV/TAP0 ... TAP7 (IV)

NAME tap0 ... tap7

SYNOPSIS --

DESCRIPTION These files refer to DECTape drives 0 to 7. Since the logical drive number can be manually set, all eight files exist even though at present there are only two physical drives.

The 256-word blocks on a standard DECTape are numbered 0 to 577. However, the system makes no assumption about this number; a block can be read or written if it exists on the tape and not otherwise. An error is returned if a transaction is attempted for a block which does not exist.

Like the RK and RF special files, addressing on the tape files is block-oriented. See the RFO section.

FILES --

SEE ALSO /dev/rf0, /dev/rk0

DIAGNOSTICS --

BUGS see /dev/rf0

OWNER ken, dmr

NAME **tty** -- console typewriter

SYNOPSIS --

DESCRIPTION **tty** (as distinct from tty0, ..., tty5) refers to the console typewriter hard-wired to the PDP-11. Most of the time it is turned off and so has little general use.

Generally, the disciplines involved in dealing with tty are similar to those for tty0 ... and the appropriate section should be consulted. The following differences are salient:

The system calls stty and gtty do not apply to this device. It cannot be placed in raw mode; on input, upper case letters are always mapped into lower case letters; a carriage return is echoed when a line-feed is typed.

The quit character is not FS (as with tty0...) but is generated by the key labelled "alt mode."

By appropriate console switch settings, it is possible to cause UNIX to come up as a single-user system with I/O on this device.

FILES --

SEE ALSO /dev/tty0...; init

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME **tty0 ... tty5** -- communications interfaces

SYNOPSIS --

DESCRIPTION These files refer to DC11 asynchronous communications interfaces. At the moment there are six of them, but the number is subject to change. Names for up to four others will be constructed by an obvious algorithm.

When one of these files is opened, it causes the process to wait until a connection is established. (In practice, however, user's programs seldom open these files; they are opened by init and become a user's standard input and output file.) The very first typewriter file open in a process becomes the control typewriter for that process. The control typewriter plays a special role in the handling quit or interrupt signals, as discussed below. The control typewriter is inherited by a child process during a fork.

A terminal associated with one of these files ordinarily operates in full-duplex mode. Characters may be typed at any time, even while output is occurring, and are only lost when the system's character input buffers become completely choked, which is very rare.

When first opened, the interface expects the terminal to use 15 odd-parity, 10-bit ASCII characters per second and to have the new-line function. Finally, the system calculates delays after sending the code for certain functions (e.g., new-line, tab) on the assumption that the terminal is a Teletype model 37. All this is merely a long way of saying that the system expects to be used by a TTY 37. However, most of these assumptions can be changed by a special system call: in particular, the expected parity can be changed; the speed, character size, and stop bits can be changed (speeds available are 134.5, 150, 300, 1200 baud; see the DC11 manual); the new-line function can be simulated by a combination of the carriage-return and line-feed functions; carriage return can be translated into new-line on input; upper case letters can be mapped into lower case letters; echoing can be turned off so the terminal operates in half duplex. See the system call stty. (Also see init for the way 300-baud terminals are detected.)

Normally, a typewriter operates in units of lines. This means that a program attempting to read will be suspended until an entire line has been typed. Also, no matter how many characters

are requested in the read call, at most one line will be returned. It is not however necessary to read a whole line at once; any number of characters may be requested in a read, even one, without losing information.

The EOT character may be used to generate an end of file from a typewriter. When an EOT is received, all the characters waiting to be read are immediately passed to the program, without waiting for a new-line. Thus if there are no characters waiting, which is to say the EOT occurred at the beginning of a line, zero characters will be passed back, and this is the standard end-of-file signal.

When the carrier signal from the dataset drops (usually because the user has hung up his terminal) any read returns with an end-of-file indication. Thus programs which read a typewriter and are sensitive to end-of-file on their inputs (which all programs should be) will terminate appropriately when hung up on.

Two characters have a special meaning when typed. The ASCII DEL character (sometimes called "rub-out") is the interrupt signal. When this character is received from a given typewriter, a search is made for all processes which have this typewriter as their control typewriter, and which have not informed the system that they wish to ignore interrupts. If there is more than one such process, one of these is selected, for practical purposes at random. Then either the process is forced to exit or a trap is simulated to an agreed-upon location in the process. See sys_intr for more information.

The ASCII character FS is the quit signal. Its treatment is identical to the interrupt signal except that unless the receiving process has made other arrangements it will not only be terminated but a core image file will be written. (See sys_quit for more information.)

During input, erase and kill processing is normally done. The character "#" erases the last character typed, except that it will not erase beyond the beginning of a line or an EOF. The character "@" kills the entire line up to the point where it was typed, but not beyond an EOF. Both these characters operate on a keystroke basis independently of any backspacing or tabbing that may have been done. Either "@" or "#" may be entered literally by preceding it by "\"; the erase or kill character remains, but the \

disappears.

It is also possible (again by sys stty) to put the typewriter into raw mode. In this mode, the program reading is wakened on each character, and when a program reads, it waits only until at least one character has been typed. In raw mode, no erase or kill processing is done; and the EOT, quit and interrupt characters are not treated specially.

Output is prosaic compared to input. It should be noted, however, that when one or more characters are written, they are actually transmitted to the terminal as soon as previously-written characters have finished typing. When a program produces characters too rapidly to be typed, as is very common, it may be suspended for a time.

Odd parity is always generated on output, except that the characters EOT and NAK have the wrong parity. Thus the 37 TTY will not hang up (EOT) or lock its keyboard (NAK) if a program accidentally prints these characters.

FILES

--

SEE ALSO

tty

DIAGNOSTICS

--

BUGS

As has been suggested, UNIX has a heavy predisposition towards 37 Teletype terminals. However, it is quite possible to use 300-baud terminals such as the GE TermiNet 300. (See init for the procedure.) The main difficulty in practice is 37-oriented delay calculations.

Terminals such as the IBM 2741 would theoretically be very desirable but there are many difficulties related to its inadequate and non-ASCII character sets (the 2741 has two, count 'em) and the inherently half-duplex nature of the terminal. It is possible to produce output on a 2741; cf type.

OWNER

ken, dmr

NAME a.out -- assembler and link editor output

SYNOPSIS --

DESCRIPTION a.out is the output file of the assembler as and the link editor ld. In both cases, a.out is executable provided there were no errors and no unresolved external references.

This file has four sections: a header, the program text, a symbol table, and relocation bits. The last two may be empty if the program was loaded with the "-s" option of ld or if the symbols and relocation have been removed by strip.

The header always contains 6 words:

- 1 a "br .+14" instruction (205(8))
- 2 The size of the program text
- 3 The size of the symbol table
- 4 The size of the relocation bits area
- 5 The size of a data area
- 6 A zero word (unused at present)

The sizes of the program, symbol table, and relocation area are in bytes but are always even. The branch instruction serves both to identify the file and to jump to the text entry point. The program text size includes the 6-word header.

The data area is used when the file is executed; the exec system call sets the program break to the sum of the text size and this data size. The data area is generated by the assembler when the location counter "." lies beyond the last assembled data, for example when the program ends with one or more constructions of the form ".=.+n"; it is preserved by the loader for the last program in a load. (Routines other than the last have the appropriate number of 0 words inserted, since there is no other provision for zero-suppression in an a.out file.)

The symbol table consists of 6-word entries. The first four contain the ASCII name of the symbol, null-padded. (In fact, the assembler generates symbols of at most 7 bytes.) The next word is a flag indicating the type of symbol. The following values are possible:

- 00 undefined symbol
- 01 absolute symbol
- 02 register symbol
- 03 relocatable symbol
- 40 undefined global symbol
- 41 absolute global symbol

43 relocatable global symbol

An undefined global corresponds to a GMAP "sym-ref" and an absolute or relocatable global to a "symdef" or absolute or relocatable value respectively. Values other than those given above may occur if the user has defined some of his own instructions.

The last word of a symbol table entry contains the value of the symbol. Its contents are not specified if the symbol is undefined.

If a.out contains no unresolved global references, header and text portions are exactly as they will appear in core when the file is executed. If the value of a word in the text portion involves a reference to an undefined global, the word is replaced by the offset in the symbol table of the appropriate symbol. (That is, possible offsets are 0, 12(10), 24(10),) Such a word will have appropriate relocation bits.

The relocation bits portion uses a variable-length encoding. There is a string of bits for each word in the text portion. The scheme has at least two bits for each word, plus possibly two more to extend the codes available; in either case the bits may be followed by a 16-bit string to represent an offset to an external symbol. The bits are packed together without regard to word boundaries. The last word is filled out with 0's on the right.

The possible relocation bit configurations are:

00

word is absolute

01

word is relocatable

10

word is a relative reference to an undefined global symbol with no offset. Currently, the word contains the offset in the symbol table of the symbol. When the symbol becomes defined, say with value x, this location will contain x-.2, where . is the location of the word.

1100xxxxxxxxxxxxxx

word is a relative reference to an external symbol with an offset. It is the same as the previous relocation type, except that the 16-bit offset is added in when the symbol

becomes defined.

1101

word is a reference to an undefined external symbol with no offset. At present the word contains the symbol table offset of the symbol. When the symbol becomes defined, the word will contain the value of the symbol.

1110xxxxxxxxxxxxxx

word is a reference to an undefined external symbol with an offset. At present, the word contains the symbol table offset of the symbol. When the symbol becomes defined, the word will contain the value of the symbol plus the given 16-bit offset.

FILES --

SEE ALSO as, ld, strip, nm, un

DIAGNOSTICS --

BUGS Soon, there will be a new type of symbol: the data area symbol. In the text, it will appear as an ordinary external reference. However, it need not be defined; this will be done by the loader. Watch this space for more details.

OWNER dmr

NAME archive (library) file format

SYNOPSIS --

DESCRIPTION The archive command ar is used to combine several files into one. Its use has three benefits: when files are combined, the file space consumed by the breakage at the end of each file (256 bytes on the average) is saved; directories are smaller and less confusing; archive files of object programs may be searched as libraries by the loader ld.

A file produced by ar has a "magic number" at the start, followed by the constituent files, each preceded by a file header. The magic number is -147(10), or 177555(8) (it was chosen to be unlikely to occur anywhere else). The header of each file is 16 bytes long:

0-7

file name, null padded on the right

8-11

Modification time of the file

12

User ID of file owner

13

file mode

14-15

file size

If the file is an odd number of bytes long, it is padded with a null byte, but the size in the header is correct.

Notice there is no provision for empty areas in an archive file.

FILES --

SEE ALSO ar, ld

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME binary punched paper tape format
SYNOPSIS --
DESCRIPTION Binary paper tape is used to pass and store arbitrary information on paper tape. The format chosen has the following features: a) no format of the data is assumed. b) check summing c) zero suppression

The format is as follows:

Between records, NULL characters are ignored. The beginning of the tape is considered between records, thus the leader is ignored.

The first non-null character specifies the type and size of the record. If the character is positive (1 to 177), the record is a data record consisting of that many characters. All but the last of these characters are data, the last being a checksum. The checksum is calculated such that the sum of the entire record is zero mod 256.

If the first character is negative (200-376) the record is a zero suppression record. It is identical to minus that number of zeros of data. One character of checksum follows this negative character. It is the positive of the negative character.

The special case of a record looking like a single zero character suppressed (377;1) causes no data transfer, but is an end-of-file indication.

FILES --
SEE ALSO lbppt, dbppt
DIAGNOSTICS --
BUGS --
OWNER ken, dmr

NAME format of core image *

SYNOPSIS --

DESCRIPTION Three conditions cause UNIX to write out the core image of an executing program: the program generates an unexpected trap (by a bus error or illegal instruction); the user sends a "quit" signal (which has not been turned off by the program); a trap is simulated by the floating point simulator. The core image is called "core" and is written in the current working directory (provided it can be; normal access controls apply). It is exactly 8192+64 bytes long. The first 8192 represent the actual contents of memory at the time of the fault; the last 64 are the contents of the system's per-user data area for this process. Only the first word of this area will be described.

When any trap which is not an I/O interrupt occurs, all the useful registers are stored on the stack. After all the registers have been stored, the contents of sp are placed in the first cell of the user area; this cell is called u.sp. Therefore, within the core image proper, there is an area which contains the following registers in the following order (increasing addresses):

(u.sp)->sc
 mq
 ac
 r5
 r4
 r3
 r2
 r1
 r0
 pc (at time of fault)
 processor status (at time of fault)

The last two are stored by the hardware. It follows that the contents of sp at the time of the fault were (u.sp) plus 22(10).

The t-bit (trap bit) in the stored status will be on when a quit caused the generation of the core image, since this bit is used in the implementation of quits.

FILES --

SEE ALSO --

DIAGNOSTICS --

11/3/71

CORE (V)

BUGS --

OWNER ken, dmr

NAME format of directories

SYNOPSIS --

DESCRIPTION A directory behaves exactly like an ordinary file, save that no user may write into a directory. The fact that a file is a directory is indicated by a bit in the flag word of its i-node entry.

Directory entries are 10 bytes long. The first word is the i-node of the file represented by the entry, if non-zero; if zero, the entry is empty.

Bytes 2-9 represent the (8-character) file name, null padded on the right. These bytes are not necessarily cleared for empty slots.

By convention, the first two entries in each directory are for ". " and "...". The first is an entry for the directory itself. The second is for the parent directory. The meaning of "... " is modified for the root directory of the master file system and for the root directories of removable file systems. In the first case, there is no parent, and in the second, the system does not permit off-device references without a mount system call. Therefore in both cases "... " has the same meaning as ". ".

FILES --

SEE ALSO file system format

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME format of file system

SYNOPSIS --

DESCRIPTION Every file system storage volume (e.g. RF disk, RK disk, DECtape reel) has a common format for certain vital information.

Every such volume is divided into a certain number of 256 word (512 byte) blocks. Blocks 0 and 1 are collectively known as the super-block for the device; they define its extent and contain an i-node map and a free-storage map. The first word contains the number of bytes in the free-storage map; it is always even. It is followed by the map. There is one bit for each block on the device; the bit is "1" if the block is free. Thus if the number of free-map bytes is n , the blocks on the device are numbered 0 through $8n-1$. The free-map count is followed by the free map itself. The bit for block k of the device is in byte $k/8$ of the map; it is offset $k \pmod{8}$ bits from the right. Notice that bits exist for the superblock and the i-list, even though they are never allocated or freed.

After the free map is a word containing the byte count for the i-node map. It too is always even. I-numbers below 41(10) are reserved for special files, and are never allocated; the first bit in the i-node free map refers to i-number 41. Therefore the byte number in the i-node map for i-node i is $(i-41)/8$. It is offset $(i-41) \pmod{8}$ bits from the right; unlike the free map, a "0" bit indicates an available i-node.

I-numbers begin at 1, and the storage for i-nodes begins at block 2. Also, i-nodes are 32 bytes long, so 16 of them fit into a block. Therefore, i-node i is located in block $(i+31)/16$ of the file system, and begins $32 \cdot ((i+31) \pmod{16})$ bytes from its start.

There is always one file system which is always mounted; in standard UNIX it resides on the RF disk. This device is also used for swapping. The swap areas are at the high addresses on the device. It would be convenient if these addresses did not appear in the free list, but in fact this is not so. Therefore a certain number of blocks at the top of the device appear in the free map, are not marked free, yet do not appear within any file. These are the blocks that show up "missing" in a check of the RF disk.

Again on the primary file system device, there

are several pieces of information following that previously discussed. They contain basically the information typed by the tm command; namely, the times spent since a cold boot in various categories, and a count of I/O errors. In particular, there are two words with the calendar time (measured since 00:00 Jan 1, 1971); two words with the time spent executing in the system; two words with the time spent waiting for I/O on the RF and RK disks; two words with the time spent executing in a user's core; one byte with the count of errors on the RF disk; and one byte with the count of errors on the RK disk. All the times are measured in sixtieths of a second.

I-node 41(10) is reserved for the root directory of the file system. No i-numbers other than this one and those from 1 to 40 (which represent special files) have a built-in meaning. Each i-node represents one file. The format of an i-node is as follows, where the left column represents the offset from the beginning of the i-node:

0-1	flags (see below)
2	number of links
3	user ID of owner
4-5	size in bytes
6-7	first indirect block or contents block
...	
20-21	eighth indirect block or contents block
22-25	creation time
26-29	modification time
30-31	unused

The flags are as follows:

100000	i-node is allocated
040000	directory
020000	file has been modified (always on)
010000	large file
000040	set user ID on execution
000020	executable
000010	read, owner
000004	write, owner
000002	read, non-owner
000001	write, non-owner

The allocated bit (flag 100000) is believed even if the i-node map says the i-node is free; thus corruption of the map may cause i-nodes to become unallocatable, but will not cause active nodes to be reused.

Byte number n of a file is accessed as follows: n is divided by 512 to find its logical block number (say b) in the file. If the file is small

(flag 010000 is 0), then b must be less than 8, and the physical block number corresponding to b is the bth entry in the address portion of the i-node.

If the file is large, b is divided by 256 to yield a number which must be less than 8 (or the file is too large for UNIX to handle). The corresponding slot in the i-node address portion gives the physical block number of an indirect block. The residue mod 256 of b is multiplied by two (to give a byte offset in the indirect block) and the word found there is the physical address of the block corresponding to b.

If block b in a file exists, it is not necessary that all blocks less than b exist. A zero block number either in the address words of the i-node or in an indirect block indicates that the corresponding block has never been allocated. Such a missing block reads as if it contained all zero words.

FILES

--

SEE ALSO

format of directories

DIAGNOSTICS

--

BUGS

Two blocks are not enough to handle the i- and free-storage maps for an RP02 disk pack, which contains around 10 million words.

OWNER

--

NAME passwd -- password file

SYNOPSIS --

DESCRIPTION passwd contains for each user the following information:

name (login name)
password
numerical user ID
default working directory
program to use as Shell

This is an ASCII file. Each field within each user's entry is separated from the next by a colon. Each user is separated from the next by a new-line. If the password field is null, no password is demanded; if the Shell field is null, the Shell itself is used.

This file, naturally, is inaccessible to anyone but the super-user.

This file resides in directory /etc.

FILES --

SEE ALSO /etc/init

DIAGNOSTICS --

BUGS --

OWNER super-user

NAME /etc/uids -- map user names to user IDs

SYNOPSIS --

DESCRIPTION This file allows programs to map user names into user numbers and vice versa. Anyone can read it. It resides in directory /etc, and should be updated along with the password file when a user is added or deleted.

The format is an ASCII name, followed by a colon, followed by a decimal ASCII user ID number.

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER dmr, ken

NAME /tmp/utmp -- user information

SYNOPSIS --

DESCRIPTION This file allows one to discover information about who is currently using UNIX. The file is binary; each entry is 16(10) bytes long. The first eight bytes contain a user's login name or are null if the table slot is unused. The low order byte of the next word contains the last character of a typewriter name (currently, '0' to '5' for /dev/tty0 to /dev/tty5). The next two words contain the user's login time. The last word is unused.

This file resides in directory /tmp.

FILES --

SEE ALSO /etc/init, which maintains the file.

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME basic -- DEC supplied BASIC

SYNOPSIS basic [file]

DESCRIPTION Basic is the standard BASIC V000 distributed as a stand alone program. The optional file argument is read before the console. See DEC-11-AJPB-D manual.

Since bas is smaller and faster, basic is not maintained on line.

FILES --

SEE ALSO bas

DIAGNOSTICS See manual

BUGS GOK

OWNER dmr

NAME

bj -- the game of black jack

SYNOPSIS

/usr/games/bj

DESCRIPTION

Black jack is a serious attempt at simulating the dealer in the game of black jack (or twenty-one) as might be found in Reno.

The following rules apply:

The bet is \$2 every hand.

A player 'natural' (black jack) pays \$3. A dealer natural loses \$2. Both dealer and player naturals is a 'push' (no money exchange).

If the dealer has an ace up, the player is allowed to make an 'insurance' bet against the chance of a dealer natural. If this bet is not taken, play resumes as normal. If the bet is taken, it is a side bet where the player wins \$2 if the dealer has a natural and loses \$1 if the dealer does not.

If the player is dealt two cards of the same value, he is allowed to 'double'. He is allowed to play two hands, each with one of these cards. (The bet is doubled also; \$2 on each hand.)

If a dealt hand has a total of ten or eleven, the player may 'double down'. He may double the bet (\$2 to \$4) and receive exactly one more card on that hand.

Under normal play, the player may 'hit' (draw a card) as long as his total is not over twenty-one. If the player 'busts' (goes over twenty-one), the dealer wins the bet.

When the player 'stands' (decides not to hit), the dealer hits until he attains a total of seventeen or more. If the dealer busts, the player wins the bet.

If both player and dealer stand, the one with the largest total wins. A tie is a push.

The machine deals and keeps score. The following questions will be asked at appropriate times. Each question is answered by **y** followed by a new line for 'yes', or just new line for 'no'.

? means 'do you want a hit?'
Insurance?

Double down?

Every time the deck is shuffled, the dealer so states and the 'action' (total bet) and 'standing' (total won or loss) is printed. To exit, hit the interrupt key (DEL) and the action and standing will be printed.

FILES	--
SEE ALSO	--
DIAGNOSTICS	--
BUGS	--
OWNER	ken

11/3/71

CAL (VI)

NAME	cal -- print calendar
SYNOPSIS	/usr/ken/cal year
DESCRIPTION	Cal will print a calendar for the given year. The year can be between 0 (really 1 BC) and 9999. For years when several calendars were in vogue in different countries, the calendar of England (and therefore her colonies) is printed.
P.S. try cal of 1752.	
FILES	--
SEE ALSO	--
DIAGNOSTICS	--
BUGS	--
OWNER	ken

NAME chess -- the game of chess
 SYNOPSIS /usr/games/chess
 DESCRIPTION Chess is an attempt at computer chess. The program 'speaks' in algebraic chess notation. The initial board configuration in this notation is as follows:

```

 8 R N B Q K B N R
 7 P P P P P P P P
 6 * * * * * * *
 5 * * * * * * *
 4 * * * * * * *
 3 * * * * * * *
 2 p p p p p p p
 1 r n b q k b n r
      a b c d e f g h
  
```

A move is specified by the 'from' co-ordinate followed by the 'to' co-ordinate. Thus the white P-K4 move would be 'e2e4'. The black P-K4 would be 'e7e5'.

The following commands are recognized by the chess program:

move

Make the move if legal. The program does not keep track of who is to play. The move is made for what ever side is specified.

move x

Make the move regardless of legality. This is a good way to either set up a desired situation or to cheat. The initial move 'e2e8x' is a winner.

mw

The program will compute and make a move for the white pieces.

m

The program will compute and make a move for the black pieces.

lab

Set the level parameters to a and b, where a and b are numbers between 0 and 9. The initial settings are 2 and 8. The first parameter increases computation time rapidly while the second parameter only increases computation exponentially. Currently move times run from 20 seconds to 10 minutes. It was hoped that these numbers would be usefully related to the program's competence.

- p The board is printed.
- u The last move is un-made. This is another good way to cheat.
- t' All the moves to date are printed.
- s The current game situation is saved on the file c.tmp.
- r The game situation on the file c.tmp is restored.
- !command
The unix command is executed by the mini-shell.

An interrupt (DEL) will pull the program out of its computation. If it is trying to make a move, the best move to date is made.

FILES	c.tmp
SEE ALSO	msh
DIAGNOSTICS	? if an illegal move is attempted, or if an unknown command is typed.
BUGS	The current version does not recognize castling, promotion and en passant. A new version is in the mill.
OWNER	ken

11/3/71

DAS (VI)

NAME das -- disassembler
SYNOPSIS --
DESCRIPTION A PDP-11 disassembler exists. Contact the author
for more information.
FILES --
SEE ALSO --
DIAGNOSTICS --
BUGS --
OWNER ken

NAME **dli** -- load DEC binary paper tapes

SYNOPSIS **dli output [input]**

DESCRIPTION **dli** will load a DEC binary paper tape into the output file. The binary format paper tape is read from the input file (/dev/ppt is default.)

FILES /dev/ppt

SEE ALSO --

DIAGNOSTICS "checksum"

BUGS --

OWNER dmr

NAME dpt -- read DEC ASCII paper tape

SYNOPSIS dpt output [input]

DESCRIPTION dpt reads the input file (/dev/ppt default) assuming the format is a DEC generated ASCII paper tape of an assembly language program. The output is a UNIX ASCII assembly program.

FILES /dev/ppt

SEE ALSO --

DIAGNOSTICS --

BUGS Almost always a hand pass is required to get a correct output.

OWNER ken, dmr

11/3/71

MOO (VI)

NAME moo -- a game
SYNOPSIS /usr/games/moo
DESCRIPTION moo is a guessing game imported from England.
FILES --
SEE ALSO --
DIAGNOSTICS --
BUGS --
OWNER ken

11/3/71

SORT (VI)

NAME sort -- sort a file
SYNOPSIS sort input output
DESCRIPTION sort will sort the input file and write the sorted file on the output file. Wide options are available on collating sequence and ignored characters.
FILES --
SEE ALSO --
DIAGNOSTICS --
BUGS --
OWNER dmr, ken

11/3/71

TTT (VI)

NAME ttt -- tic-tac-toe
SYNOPSIS /usr/games/ttt
DESCRIPTION ttt is the X's and O's ~~the is~~ popular in 1st
grade. This is a learning program that never
makes the same mistake twice.
game
FILES ttt.k -- old mistakes
SEE ALSO --
DIAGNOSTICS --
BUGS --
OWNER ken

11/3/71

/ETC/AS2 (VII)

NAME as2 -- assembler pass 2

SYNOPSIS --

DESCRIPTION as2 is invoked by the assembler as to perform its second pass.

FILES see as

SEE ALSO as

DIAGNOSTICS see as

BUGS --

OWNER dmr

11/3/71

/ETC/BA (VII)

NAME **ba** -- B assembler

SYNOPSIS /etc/ba name

DESCRIPTION **ba** is invoked by the **B** command in order to turn
 the **B** intermediate code into assembly language.

FILES name.i (input), name.s (output)

SEE ALSO **b** command, /etc/bc

DIAGNOSTICS --

BUGS At the moment, the **b** command is defunct, and **ba**
 is invoked via a command file.

OWNER ken

NAME **bc** -- B compiler

SYNOPSIS /etc/bc name.b name.i

DESCRIPTION **bc** is the B compiler proper; it turns B source into intermediate code. It is invoked from the b command.

FILES name.b (input), name.i (intermediate output)

SEE ALSO b (command), /etc/ba

DIAGNOSTICS --

BUGS The b command is defunct at the moment; bc is called from a command file.

OWNER ken

NAME **bilib** -- B interpreter library

SYNOPSIS --

DESCRIPTION **bilib** is the library of B runtime operators. It is searched during the loading of a B-compiled program.

Standard B subroutines are contained in /etc/libb.a.

FILES --

SEE ALSO b (command); ar, ld

DIAGNOSTICS --

BUGS The following assignment binary operators are missing: b102 (=|), b103 (=&), b104 (==), b105 (!=), b106 (=<=), b107 (=<), b110 (=>=), b111 (=>), b112 (=>>), b113 (=<<), b120 (=/).

OWNER ken, dmr

NAME **bos, maki, rom, vcboot, msys, et al**

SYNOPSIS --

DESCRIPTION On the RF disk, the highest 16K words are reserved for stand-alone programs. These 16K words are allocated as follows:

bos	(1K)
Warm UNIX	(6K)
Cold UNIX	(6K)
unassigned	(3K)

The UNIX read only memory (ROM) is home cut with 2 programs of 16 words each. The first (address 173700) reads bos from the RF disk into core location 54000 and transfers to 54000. The other ROM program (address 173740) reads a DECtape sitting in the end-zone on drive 0 into core location 0 and transfers to 0. This latter operation is compatible with part of DEC's standard ROM. The disassembled code for the UNIX ROM follows:

173700:	mov	\$177472,r0	12700;177472
	mov	\$3,-(r0)	12740;3
	mov	\$140000,-(r0)	12740;140000
	mov	\$54000,-(r0)	12740;54000
	mov	\$-2000,-(r0)	12740;176000
	mov	\$5,-(r0)	12740;5
	tstb	(r0)	105710
	bge	.-2	2376
	jmp	*\$54000	137;54000
173740:	mov	\$177350,r0	12700;177350
	clr	-(r0)	5040
	mov	r0,-(r0)	10040
	mov	\$3,-(r0)	12740;3
	tstb	(r0)	105710
	bge	.-2	2376
	tst	*\$177350	5737;177350
	bne	.	1377
	movb	\$5,(r0)	112710;5
	tstb	(r0)	105710
	bge	.-2	2376
	clr	pc	5007

The program bos (Bootstrap Operating System) examines the console switchs and executes one of several internal programs depending on the setting. If no setting is recognizable, bos loops waiting for a recognizable setting. The following settings are currently recognized:

173700
 73700 Will read Warm UNIX from the RF into core location 0 and transfer to 400.

- 1 Will read Cold UNIX from the RF into core location 0 and transfer to 400.
- 2 Will read the unassigned 3K program into core location 0 and transfer to 400.
- 10 Will dump 12K words of memory from core location 0 onto DECTape drive 7.
- 0 Will load a standard UNIX binary paper tape into core location 0 and transfer to 0.
- 57500 Will load the standard DEC absolute and binary loaders and transfer to 57500.

Thus we come to the UNIX warm boot procedure: put 173700 into the switches, push load address and then push start. The alternate switch setting of 73700 that will load warm UNIX is used as a signal to bring up a single user system for special purposes. See /etc/init.

Cold boots can be accomplished with the Cold UNIX program, but they're not. Thus the Cold UNIX slot on the RF may have any program desired. This slot is, however, used during a cold boot. Mount the UNIX INIT DECTape on drive 0 positioned in the end-zone. Put 173740 into the switches. Push load address. Put 1 into the switches. Push start. This reads a program called vcboot from the tape into core location 0 and transfers to it. vcboot then reads 16K words from the DECTape (blocks 1-32) and copies the data to the highest 16K words of the RF. Thus this initializes the read-only part of the RF. vcboot then reads in bos and executes it. bos then reads in Cold UNIX and executes that. Cold UNIX halts for a last chance before it completely initializes the RF file system. Push continue, and Cold UNIX will initialize the RF. It then sets into execution a user program that reads the DECTape for initialization files starting from block 33. When this is done, the program executes /etc/init which should have been on the tape.

The INIT tape is made by the program maki running under UNIX. maki writes vcboot on block 0 of /dev/tap7. It then copies the RF 16K words (using /dev/rf0) onto blocks 1 thru 32. It has internally a list of files to be copied from block 33 on. This list follows:

/etc/init
/bin/chmod

```
/bin/chown  
/bin/cp  
/bin/ln  
/bin/ls  
/bin/mkdir  
/bin/mv  
/bin/rm  
/bin/rmdir  
/bin/sh  
/bin/stat  
/bin/tap
```

Thus this is the set of programs available after a cold boot. /etc/init and /bin/sh are mandatory. /bin/tap and /bin/mkdir are used to load up the file system. The rest of the programs are frosting. As soon as possible, an sdate should be done.

The last link in this incestuous daisy chain is the program msys.

msys char file

will copy the file file onto the RF read only slot specified by the character char. Char is taken from the following set:

<u>b</u>	<u>bos</u>
<u>u</u>	Warm UNIX
<u>1</u>	Cold UNIX
<u>2</u>	unassigned

Due to their rarity of use, maki and msys are maintained off line and must be reassembled before used.

FILES	/dev/rf0, /dev/tapn
SEE ALSO	/etc/init, /bin/tap, /bin/sh, /bin/mkdir, bppt format
DIAGNOSTICS	--
BUGS	The files /bin/mount, /bin/sdate, and /bin/date should be included in the initialization list of <u>maki</u> .
OWNER	ken

11/3/71

/ETC/BRT1, BRT2 (VII)

NAME brt1, brt2 -- B runtime routines

SYNOPSIS --

DESCRIPTION The first of these routines must be loaded first in an executable B program; the second must be loaded last, after all other routines. They are not in /etc/bilib only because having them separate is the easiest way to assure the order of loading.

FILES --

SEE ALSO b command, bilib

DIAGNOSTICS --

BUGS --

OWNER ken

11/3/71

/ETC/F1, F2, F3, F4 (VII)

NAME f1, f2, f3, f4 -- Fortran compiler

SYNOPSIS --

DESCRIPTION These programs represent the four phases of a Fortran compilation:

f1: specification statements
f2: common and equivalence allocation
f3: executable statements
f4: cleanup

Each exec's the next; the first is called by the for command.

FILES f.tmp1, f.tmp2, f.tmp3

SEE ALSO for

DIAGNOSTICS --

BUGS Besides the fact that there is a good deal of the Fortran language missing, there is no for command; Fortran is invoked via a command file.

OWNER ken, dmr

11/3/71

/ETC/GLOB (VII)

NAME	glob -- global
SYNOPSIS	--
DESCRIPTION	<u>glob</u> is used to expand arguments to the shell containing "*" or "?". It is passed the argument list containing the metacharacters; <u>glob</u> expands the list and calls the command itself.
FILES	--
SEE ALSO	sh
DIAGNOSTICS	"No match", "no command"
BUGS	<u>glob</u> will only load a command from /bin. Also if any "*" or "?" argument fails to generate matches, "No match" is typed and the command is not executed.
OWNER	dmr

NAME init -- process initialization

SYNOPSIS --

DESCRIPTION init is invoked inside UNIX as the last step in the boot procedure. It first carries out several housekeeping duties: it must change the modes of the tape files and the RK disk file to 17, because if the system crashed while a tap or rk command was in progress, these files would be inaccessible; it also truncates the file /tmp/utmp, which contains a list of UNIX users, again as a recovery measure in case of a crash. Directory usr is assigned via sys mount as resident on the RK disk.

init then forks several times so as to create one process for each typewriter channel on which a user may log in. Each process changes the mode of its typewriter to 15 (read/write owner, write-only non-owner; this guards against random users stealing input) and the owner to the super-user. Then the typewriter is opened for reading and writing. Since these opens are for the first files open in the process, they receive the file descriptors 0 and 1, the standard input and output file descriptors. It is likely that no one is dialled in when the read open takes place; therefore the process waits until someone calls. At this point, init types its "login:" message and reads the response, which is looked up in the password file. The password file contains each user's name, password, numerical user ID, default working directory, and default shell. If the lookup is successful and the user can supply his password, the owner of the typewriter is changed to the appropriate user ID. An entry is made in /tmp/utmp for this user to maintain an up-to-date list of users. Then the user ID of the process is changed appropriately, the current directory is set, and the appropriate program to be used as the Shell is executed.

At some point the process will terminate, either because the login was successful but the user has now logged out, or because the login was unsuccessful. The parent routine of all the children of init has meanwhile been waiting for such an event. When return takes place from the sys wait, init simply forks again, and the child process again awaits a user.

There is a fine point involved in reading the login message. UNIX is presently set up to handle automatically two types of terminals: 150 baud, full duplex terminals with the line-feed

function (typically, the Model 37 Teletype terminal), and 300 baud, full duplex terminals with only the line-space function (typically the GE TermiNet terminal). The latter type identifies itself by sending a line-break (long space) signal at login time. Therefore, if a null character is received during reading of the login line, the typewriter mode is set to accommodate this terminal and the "login:" message is typed again (because it was garbled the first time).

Init, upon first entry, checks the switches for 73700. If this combination is set, init will open /dev/tty as standard input and output and directly execute /bin/sh. In this manner, UNIX can be brought up with a minimum of hardware and software.

FILES /tmp/utmp, /dev/tty0 ... /dev/ttyn

SEE ALSO sh

DIAGNOSTICS "No directory", "No shell". There are also some halts if basic I/O files cannot be found in /dev.

BUGS --

OWNER ken, dmr

NAME kbd -- keyboard map

SYNOPSIS cat /etc/kbd

DESCRIPTION kbd contains a map to the keyboard for model 37 Teletype terminals with the extended character set feature. If kbd is printed on such a terminal, the following will appear:

```
<[1234567890-_] ^\ >qwertyuiop@ asdfghjkl;: zxcvbnm,./
<◊1234567890-¬δ∫Υ >                         ;:                ,./

<{!"#$%&'() =_}~| >QWERTYUIOP` ASDFGHJKL+* ZXCVBNM,.?
< !"#$_%&'() =¬ >γΔΛΣθσφτθΠ αεδΦΓΨπρλ+* Ωξωψβημ,.?
```

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER jfo

11/3/71

/ETC/LIBA.A (VII)

NAME liba.a -- assembly language library

SYNOPSIS --

DESCRIPTION This library is the standard location for assembly-language subroutines of general use. A section of this manual is devoted to its contents.

This library is searched when the link editor ld encounters the "-l" argument.

FILES --

SEE ALSO ld; library manual

DIAGNOSTICS --

BUGS --

OWNER dmr, ken

NAME libb.a -- B library

SYNOPSIS --

DESCRIPTION This library contains all B-callable subroutines of general utility. Its contents are detailed in the library section of the B manual. At present its contents are:

```
char
getchr
putchr
exit
printf
seek
setuid
stat
time
unlink
wait
lchar
chdir
chmod
chown
close
creat
execl
execv
fork
fstat
getuid
intr
link
makdir
open
read
write
ctime
```

FILES --

SEE ALSO b

DIAGNOSTICS --

BUGS --

OWNER ken, dmr

NAME /etc/libf.a -- Fortran library

SYNOPSIS --

DESCRIPTION This library contains all the Fortran runtime routines. Many are missing.

FILES --

SEE ALSO f1, f2, f3, f4

DIAGNOSTICS --

BUGS Will be renamed, and libf.a reserved for subroutines and functions.

OWNER ken, dmr

NAME logging in and logging out

SYNOPSIS --

DESCRIPTION UNIX must be called from an appropriate terminal. The two general classes of terminals which UNIX supports are typified by the 37 Teletype on the one hand and the GE TermiNet 300 and Memorex 1240 on the other. The principal difference is the baud rate (150 vs. 300) and the treatment of the carriage return character. Most terminals operating at 150, 300, or 1200 baud using the ASCII character set either work (more or less) at the moment or can be used by special arrangement. In particular, special arrangement is necessary for terminals which do not generate lower-case ASCII characters.

It is also necessary to have a valid UNIX user ID and (if desired) password. These may be obtained, together with the telephone number, from the system administrators.

The same telephone number serves terminals operating at both the standard speeds. When a connection is established via a 150-baud terminal (e.g. TTY 37) UNIX types out "login:"; you respond with your user name, and, if a mask is typed, with a password. If the login was successful, the "@" character is typed by the Shell to indicate login is complete and commands may be issued. A message of the day may be typed if there are any announcements. Also, if there is a file called "mailbox", you are notified that someone has sent you mail. (See the mail command.)

From a 300-baud terminal, the procedure is slightly different. Such terminals often have a full-duplex switch, which should be turned on (or conversely, half-duplex should be turned off). When a connection with UNIX is established, a few garbage characters are typed (these are the "login:" message at the wrong speed). You should depress the "break" key; this is a speed-independent signal to UNIX that a 300-baud terminal is in use. It will type "login:" (at the correct speed this time) and from then on the procedure is the same as described above.

Logging out is simple by comparison (in fact, sometimes too simple). Simply generate an end-of-file at Shell level by using the EOT character; the "login:" message will appear again to indicate that you may log in again.

It is also possible to log out simply by hanging up the terminal; this simulates an end-of-file on the typewriter.

FILES	--
SEE ALSO	<u>init</u>
DIAGNOSTICS	--
BUGS	Hanging up on programs which never read the type-writer or which ignore end-of-files is very dangerous; in the worst cases, the programs can only be halted by restarting the system.
OWNER	ken, dmr

11/3/71

/ETC/MSH (VII)

NAME	msh -- mini-shell
SYNOPSIS	--
DESCRIPTION	<p><u>msh</u> is a heavily simplified version of the Shell. It reads one line from the standard input file, interprets it as a command, and calls the command.</p> <p>The mini-shell supports few of the advanced features of the Shell; none of the following characters is special:</p> <p style="text-align: center;">> < \$ \ ; &</p> <p>However, "*" and "?" are recognized and <u>glob</u> is called. The main use of <u>msh</u> is to provide a command-executing facility for various interactive sub-systems.</p>
FILES	--
SEE ALSO	sh, glob
DIAGNOSTICS	"?"
BUGS	--
OWNER	ken, dmr

NAME **suftab** -- suffix table

SYNOPSIS --

DESCRIPTION **suftab** is a table of suffixes used to guide hyphenation in roff. Its first 12 words are not used (see a.out format.) Its next 26 words point to the beginning of the subtables for each of the 26 initial letters of a suffix. The first entry for each suffix is a count of the number of bytes in the suffix. The second byte of each entry is a flag indicating the type of suffix. The suffix itself follows; the high bits of each letter indicate where the hyphens come. The table for each initial suffix letter ends with a zero count byte.

FILES --

SEE ALSO roff

DIAGNOSTICS --

BUGS --

OWNER jfo, dmr, ken

NAME tabs -- tab stop set

SYNOPSIS cat /etc/tabs

DESCRIPTION When printed on a suitable terminal, this file will set tab stops at columns 8, 16, 24, 32, Suitable terminals include the Teletype model 37 and the GE TermiNet 300.

Since UNIX times delays assuming tabs set every 8, this has become a defacto 'standard.'

FILES --

SEE ALSO --

DIAGNOSTICS --

BUGS --

OWNER ken